King's Russia Institute
Dissertation Cover Sheet and Feedback Form

Please complete all sections of this form and securely attach it to the front of your dissertation that you are submitting for assessment.
DO NOT PUT YOUR NAME ANYWHERE ON YOUR WORK, BUT MAKE SURE THAT YOUR CANDIDATE NUMBER APPEARS ON EVERY PAGE.
Failure to attach this form as required may result in your work not being accepted for assessment.

Student's candidate number:	AF13171
Programme:	**MSc Russian & Eurasian Politics and Economics**
Dissertation title:	Quantifying Persuasion: Mapping Hidden Influence in the Podcast Election
Essay word count*:	14110

First marker's comments:

Second marker's comments:

***The word count, which should preferably be calculated electronically, must be stated accurately above.** Please refer to the Student Handbook or module syllabus for word limit regulations.

By submitting an assignment under this cover sheet, you are signifying your assent to the following declaration.
DECLARATION BY STUDENT
This assignment is entirely my own work. Quotations from secondary literature are indicated by the use of inverted commas around ALL such quotations AND by reference in the text or notes to the author(s) concerned. ALL primary and secondary literature used in this piece of work is indicated in the bibliography placed at the end, and dependence upon ANY source used is indicated at the appropriate point in the text. I confirm that no sources have been used other than those stated. I UNDERSTAND WHAT IS MEANT BY PLAGIARISM AND HAVE SIGNED THE DECLARATION CONCERNING THE AVOIDANCE OF PLAGIARISM. I UNDERSTAND THAT PLAGIARISM IS A SERIOUS EXAMINATIONS OFFENCE THAT MAY RESULT IN DISCIPLINARY ACTION BEING TAKEN.

Quantifying Persuasion

Mapping Hidden Influence in the Podcast Election

AF13171

September 2025

MSc Russian and Eurasian Politics and Economics

Russia Institute

King's College London

Supervisor: Dr. Gulnaz Sharafutdinova

Abstract

This dissertation examines the Tenet Media case, a 2024 U.S. Department of Justice indictment that exposed a Russian state-backed influence operation financing American and Canadian political podcasters. While conventional approaches to disinformation research prioritise attribution, this study argues that such efforts fail in environments where foreign-funded content is structurally indistinguishable from domestic discourse. Instead, it advances an attribution-agnostic framework of rhetorical fingerprinting, treating manipulation as legible through the effects of discourse rather than its origins.

Drawing on a corpus of 6,437 political podcast episodes (including 405 Tenet Media transcripts), the study operationalises a six-dimensional computational framework (causal reasoning, emotional activation, persuasive techniques, thematic saturation, sentiment calibration, and topical clustering) integrating symbolic and transformer-based methods. Findings demonstrate that Tenet's rhetorical profile aligned by 94–96% with U.S. right-wing podcasts, employing near-identical causal structures, emotional arcs, and persuasive strategies. Tenet rarely mentioned Russia, concentrated on U.S. political figures, and operated at mid-level "professionalised" intensity, embedding itself seamlessly within domestic discourse. Classification models performed no better than chance in distinguishing Tenet from right-wing podcasts, confirming that attribution-based detection is mathematically undermined by baseline convergence.

By reframing information operations as environmental manipulations rather than anomalies traceable to foreign actors, the study proposes rhetorical fingerprinting as a scalable method for monitoring discursive health. This approach enables measurement of influence intensity, stylistic alignment, and narrative saturation across ecosystems without adjudicating truth claims or censoring speech. The findings underscore a broader theoretical shift: influence must be studied as an environmental science, where resilience depends not on isolating adversaries but on continuously monitoring rhetorical conditions that adversaries exploit.

Table of Contents

Abstract *3*
CHAPTER 1: INTRODUCTION 6
1.1 When Attribution Becomes Irrelevant: The Tenet Disruption *6*
1.2 The Archaeology of Manipulation: From Active Measures to Environmental Control *6*
1.3 The Intimacy Weapon: Podcasts as Vectors for Parasocial Exploitation *7*
1.4 The Forensic Opportunity: Beyond Attribution Toward Environmental Measurement *7*
Research Questions *8*
CHAPTER 2: LITERATURE REVIEW 8
2.1 Podcasts as Information Warfare Vectors *8*
2.2 The Attribution Paradox in Russian Operations Research *10*
2.3 Evolution of Russian Information Operations *12*
2.4 Computational Approaches to Influence Detection *12*
2.5 Synthesis: Toward Rhetorical Fingerprinting *14*
CHAPTER 3: METHODOLOGY 16
3.1 Operationalising Rhetorical Fingerprinting *16*
3.2 The Six-Dimensional Framework *17*
3.3 Corpus as Environmental Sample *18*
3.4 Three-Stage Processing Pipeline *19*
3.5 Blame Attribution: The Six Dimensions in Action *25*
3.6 Tenet Media Temporal Limitation *27*
3.7 A Note on Scalability: Segment-Level Search Architecture and Metadata Structure *27*
CHAPTER 4: FINDINGS 28
4.1 The Baseline Discovery *28*
Analysis of 6,437 podcast episodes reveals a striking finding: Russian-funded Tenet Media operates at an influence intensity statistically indistinguishable from American political podcasts. Bootstrap resampling (1,000 iterations) established baseline means with 95% confidence intervals: Tenet Media scored 5.454 [5.218, 5.700], U.S. right-wing podcasts 5.684 [5.630, 5.737], and U.S. left-wing podcasts 4.782 [4.720, 4.844]. The overlap between Tenet and the right-wing interval ($p > 0.05$) demonstrates that the foreign-funded cluster falls squarely within the natural range of domestic discourse. *28*
4.2 The Influence Hierarchy and Individual Fingerprints *29*
4.3 Structural Convergence: The 94% Alignment *31*
4.4 Thematic Fingerprinting: What They Talk About *32*
4.5 Political Figure Analysis: Who Gets Discussed *33*
4.6 Persuasive Technique Deployment *34*
4.7 Persuasive Arcs and Emotional Dynamics *35*
4.8 Blame Attribution Networks *38*
4.9 The Russia Paradox: Strategic Minimisation *39*

4.10 The Baseline Problem: Mathematical Undetectability *40*
4.11 Multidimensional Analysis Showcase: The Strategic Minimisation Pattern *41*
4.12 Temporal Limitations and Baseline Comparison *42*
4.13 Summary of Key Findings *43*
CHAPTER 5: STRUCTURAL INFLUENCE: FINDINGS ON INTENSITY, STYLISTICS, AND NARRATIVE ALIGNMENT 43
5.1 From Attribution to Expansion *43*
5.2 Visible Dissent in the Digital Age: Why Intensity Matters *45*
5.3 Stylistic Indicators and Measurement *45*
5.4 Technical and Stylistic Advances *46*
5.5 Stylistic Integration and Environmental Drift *46*
5.6 Implications for Democratic Resilience *47*
CHAPTER 6: THE ARCHITECTURE OF RESILIENCE: MONITORING RHETORICAL ENVIRONMENTS 47
6.1 Beyond Attribution: A New Framework for Democratic Defence *47*
6.2 Technical Infrastructure: Scalable Monitoring Architectures *48*
6.3 Future Directions: Building Infrastructures of Resilience *48*
6.4 Stakeholder Intelligence and Applications *49*
6.5 Safeguards for Democracy *50*
6.6 Conclusion: The Environmental Turn *51*

Table of Figures:

Figure 3.1: Data Processing Pipeline	Pg. 17
Figure 3.2: 6-Dimensional Framework	Pg. 24
Figure 4.1: Distribution of Influence Intensity	Pg. 28
Figure 4.2: Influence Intensity by Show	Pg. 30
Figure 4.3: Emotional Sequencing Convergence	Pg. 32
Figure 4.4: Thematic Convergence	Pg. 33
Figure 4.5: The Russia Paradox	Pg. 40
Figure 5.1: The Rhetorical Phalanx	Pg. 44

Table of Tables:

Table 3.1: Corpus Composition Summary	Pg. 19
Table 3.2 Empirical Weighting Validation	Pg. 24
Table 4.1: Complete Show Rankings	Pg. 31
Table 4.2: Political Figure Mention Rates	Pg. 34
Table 4.3: Persuasive Technique Usage	Pg. 35
Table 4.4: Blame Targets	Pg. 39
Table 4.6: Multidimensional Trump Analysis	Pg. 42
Table 4.7: Multidimensional Biden Analysis	Pg. 42

CHAPTER 1: Introduction

1.1 When Attribution Becomes Irrelevant: The Tenet Disruption

On September 4, 2024, the U.S. Department of Justice unsealed an indictment charging two Russian nationals, Kostiantyn Kalashnikov and Elena Afanasyeva, both employees of the state-controlled network RT, with orchestrating a covert influence operation through a Tennessee-based front company known as Tenet Media (DOJ, 2024; Marchman & Mehrotra, 2024; McCausland, 2024). Nearly $10 million was covertly funnelled through shell corporations to finance U.S. and Canadian political commentators, who produced almost 2,000 videos across YouTube, TikTok, Instagram, X, and Rumble, generating over 16 million YouTube views (DOJ, 2024).

The indictment referred only to "Commentator-1" and "Commentator-2," describing their subscriber bases but not naming them (DOJ, 2024). Subsequent reporting identified the influencers as Tim Pool, Benny Johnson, Dave Rubin, Lauren Southern, Matt Christiansen, and Tayler Hansen, among others, and that these creators, who collectively reached audiences exceeding 20 million, received payments sometimes disguised as corporate investments allegedly reaching up to $100,000 per video (Marchman & Mehrotra, 2024). When the operation was exposed, several insisted they were "victims" who retained editorial control (McCausland, 2024).

Yet this claim misses the point. As the WIRED transcript archive compiled by Marchman and Mehrotra (2024) demonstrates, the content consistently aligned with Kremlin objectives: sowing scepticism of U.S. support for Ukraine, delegitimising NATO, and amplifying divisive domestic grievances. The Tenet case exemplifies reflexive control, a Soviet doctrine of shaping adversaries' choices by manipulating the information environment while preserving the illusion of autonomy (De Goeij, 2023). Its power lies in deniability: how does one prove influence when those influenced sincerely believe the words are their own?

1.2 The Archaeology of Manipulation: From Active Measures to Environmental Control

The Tenet operation is not an aberration, but the latest stage in a lineage of active measures. As Rid (2020) observes, disinformation operations are "covert operations designed to achieve overt influence, carefully hidden yet visible in plain sight" (p. 9).

During the 2016 U.S. election, public attention focused on the Internet Research Agency (IRA)'s bots and memes, the visible surface layer. Yet as the Mueller Report (2019) documented, the more profound disruption came from hack-and-leak operations (Guccifer 2.0, DCLeaks), which reframed political narratives while analysts chased troll activity. Subsequent research showed how the IRA sustained and expanded operations across platforms. Howard et al. (2018) found that the IRA targeted cultural identity divides and intensified U.S. polarisation, particularly via Instagram. The Senate-commissioned report by DiResta et al. (2019) detailed tactical tropes such as memetics, voter suppression, and parasocial trust-building that became the standard playbook.

These campaigns exemplify what Starbird, Arif, and Wilson (2019) call disinformation as collaborative work; strategic operations do not only rely on paid actors but also enlist unwitting participants, weaving into existing online communities and leveraging authentic grievances. This framework helps explain Tenet's design: it did not invent false narratives but amplified authentic domestic voices, aligning them with Russian strategic goals. This practice, which might be called narrative arbitrage, turns preexisting viewpoints into vehicles of foreign influence. Tucker (2018) notes that propaganda may consist entirely of accurate information deployed to disparage rivals.

1.3 The Intimacy Weapon: Podcasts as Vectors for Parasocial Exploitation

Among the platforms Tenet used, podcasts proved most potent. Pew Research (2023) reports that over 40% of Americans listen monthly, and 87% rate podcasts as "mostly accurate," a higher credibility rating than broadcast or print. Their persuasive power derives from parasocial intimacy: the "quiet voice nestled cosily, almost out of sight, and chosen just for you" (Sienkiewicz & Jaramillo, 2019).

This intimacy makes podcasts ideal channels for long-form influence. Unlike tweets or op-eds, hours of conversation dissolve argument into "common sense," where speculation about civil war or corruption registers less as a claim than as confidence from a trusted voice. The architecture compounds the risk: podcasts are distributed through decentralised RSS feeds with little centralised moderation or transparency (Wirtschafter, 2021).

McLuhan's dictum that "the medium is the message" (1964) underscores why this matters: it is not only the content of podcasts but the medium itself but its intimacy, decentralisation, and credibility that reshapes political legitimacy. By 2024, podcasts had become central to campaign strategy, with candidates prioritising podcast appearances over mainstream media. Trump's Joe Rogan interview, which drew over 50 million views, cemented what came to be called "the podcast election" (Brimson, N. H. & Lemon, L. L., 2023). For foreign operators, exploiting this medium means exploiting the very channels through which democratic legitimacy is now constructed.

1.4 The Forensic Opportunity: Beyond Attribution Toward Environmental Measurement

The Tenet case provides a rare artefact: confirmed foreign-funded content embedded in authentic American discourse, operating undetected for months. Unlike traditional propaganda detection models focused on attribution, this case allows for methodological innovation.

This dissertation conducts a comparative rhetorical analysis of Tenet podcasts against domestic U.S. shows, using the transcript corpus compiled by Marchman and Mehrotra (2024) as primary data. Rather than tracing origins, it follows a detection-over-attribution framework, focusing on rhetorical structures of influence that persist regardless of source.

This aligns with recent scholarship (Yang et al., 2024), which shows that Russian outlets like RT flexibly lead on geopolitics but follow U.S. outlets on divisive domestic issues, adapting tactics to serve strategic interests. It also resonates with Starbird et al.'s (2019) argument that

disinformation is best understood as environmental manipulation: altering salience and framing rather than fabricating lies.

Research Questions

RQ1: What rhetorical structures recur in long-form political podcast discourse, and how can they be fingerprinted computationally for intensity and influence monitoring?

RQ2: What rhetorical fingerprints distinguish Russian-funded podcast content from domestic American discourse?

RQ3: What do these patterns reveal about the broader vulnerability of American political communication infrastructure?

To answer these, the study employs a six-dimensional framework: causal reasoning patterns, emotional tonality, blame attribution, persuasive density, thematic clustering, and influence intensity. These features are extracted through a hybrid of regex-based symbolic methods and transformer-based inference, following precedents in propaganda detection (Bassi et al., 2024; Da San Martino et al., 2019).

The goal is not surveillance but systems-level monitoring, modelled on large-scale Elasticsearch-based monitoring platforms, to enable real-time fingerprinting of long-form narratives and adaptive response. By developing such tools, this research seeks to equip campaigns, officials, researchers, and institutions to better navigate the shifting rhetorical environment that foreign influence operations have learned to exploit.

Chapter 2: Literature Review

2.1 Podcasts as Information Warfare Vectors

The podcast medium has emerged as a particularly potent vector for information warfare, one that is strikingly susceptible to manipulation yet remarkably difficult to monitor. Unlike traditional media, which typically adhere to institutional editorial oversight, or social media, which is often subject to content moderation, podcast audio is distributed through decentralised RSS feeds and lacks central regulation. Usually, content must be transcribed from multiple sources before it's analysable, a laborious process that impairs timely oversight. As Valerie Wirtschafter (2023) puts it, "podcasts have generally offered a space … 'you can say whatever you want'," reflecting the medium's regulatory vacuum. This decentralisation creates ample space for both domestic and foreign actors to insert misinformation without a meaningful counterbalance.

Wirtschafter's study (2023) offers one of the most comprehensive assessments to date. Analysing 36,603 episodes from 79 popular political podcasts, the research found that nearly 70 % of podcasters aired at least one unsubstantiated or false claim, and about one in every 20 episodes featured such content. These findings reinforce the notion of podcasts as dark matter of the information ecosystem, ubiquitous yet poorly illuminated. Perhaps more troubling, her work reveals that false claims spiked dramatically after the 2020 U.S. presidential election, particularly in conservative-leaning podcasts. In the lead-up to January 6, 2021, episodes containing election fraud narratives jumped by nearly 600%, comprising over 28% of all episodes in that period (Wirtschafter, 2023). This behaviour underscores the role of podcasts as "narrative incubators": spaces where politicised misinformation gains traction before crossing into mainstream discourse, aligning with Brandt et al.'s (2022) findings on cross-media amplification.

Podcasting's intimacy and parasocial appeal further amplify its potential as an influence vector. Listeners often form deep trust bonds with hosts (Sharon & John, 2024; Sienkiewicz & Jaramillo, 2019), leading them to accept content at face value. Pew Research (2023) reports that 87% of podcast news consumers expect content to be "mostly accurate", a trust level considerably higher than for traditional media. Brandt et al. (2022) concretely demonstrate this dynamic in U.S. political discourse, showing how disinformation about bioweapons labs in Ukraine spread through political podcasts reaching millions before institutional rebuttals, and then escalated via mainstream TV.

On the psychological front, parasocial dynamics intensify susceptibility. Schramm et al. (2024), in their review of 281 studies on parasocial interactions, emphasise that auditory media such as radio and podcasts are especially fertile environments for parasocial bonding. They note that the human voice itself is a powerful trigger of familiarity and trust, and that many podcast listeners report tuning in primarily for their attachment to the host, even when they have little prior interest in the episode's topic. These findings suggest that podcast hosts occupy a role similar to trusted radio presenters, with parasocial relationships driving both sustained engagement and heightened receptivity to the content itself. Girard (2024) underscores this dynamic in her analysis of women-hosted podcasts that reframed the January 6th Capitol attack through domestic and feminine-coded narratives, giving rise to "intimate publics" that facilitated radicalisation through socially acceptable channels. Shekhovtsov (2017) shows how Russian media platforms such as RT deliberately cultivate Western far-right and conspiracy figures, offering them visibility and legitimacy as part of a broader ecosystem of anti-systemic voices. Because these actors already share narratives critical of liberal democracy, their amplification requires little overt direction, making them effective vehicles for laundering Kremlin-aligned messages into domestic discourse.

Beyond foreign influence, domestic investment has accelerated podcast growth by reinforcing the same structural incentives. Prentis Jacobs et al. (2024) document that U.S. audiences now exceed 100 million weekly listeners, with advertising revenues projected to reach nearly $2.6 billion by 2026. This influx reflects audience expansion and changing industry perceptions: where advertisers once worried about brand safety in a largely unregulated environment, evidence now shows that listeners are receptive to podcast advertising and perceive sponsored content as credible and engaging. These dynamics have encouraged major U.S. brands to treat podcasting as a premium growth channel, pouring investment into the medium at a time when returns on traditional digital advertising are plateauing. Thus, the same deregulation that makes

podcasts vulnerable to manipulation also makes them commercially attractive, drawing in both political actors and mainstream capital.

2.2 The Attribution Paradox in Russian Operations Research

A defining challenge in identifying and mitigating foreign influence operations is the problem of attribution. Scholars disagree not only on whether attribution is possible, but also on whether it is even the most critical question to ask. At the heart of this debate lies an epistemological tension: is it more valuable to know what manipulation does, or who is behind it? The first orientation emphasises effects, such as agenda-setting, polarisation, or discourse shifts; the second privileges origins, seeking to trace responsibility back to identifiable actors, states, or proxies. This tension has divided the literature into two broad categories: structural detection approaches, which identify influence through recurring fingerprints and measurable impacts, and critiques of attribution, which argue that tracing responsibility is methodologically flawed or strategically irrelevant.

Structural Detection Approaches

Research in this category advances the view that attribution is unnecessary so long as manipulation can be detected through its effects. Yang, McCabe, and Hindman (2024) exemplify this orientation in their analysis of 4.7 million news articles and 22.6 billion Facebook interactions across U.S. media. Using structural topic modelling, they show that Russian outlets RT and Sputnik consistently led U.S. media coverage on geopolitical issues such as international trade and security, while following centrist and far-right outlets on divisive domestic topics. Their findings demonstrate that propaganda's influence can be detected through agenda-setting patterns even without tracing the identities of individual operators.

A similar model is developed in DiResta et al. (2019), whose Tactics and Tropes of the Internet Research Agency report maps the recurring rhetorical strategies of Russian troll campaigns. These included identity polarisation, distrust of institutions, and memetic warfare, which appeared consistently across millions of Facebook, Instagram, and Twitter posts. Such recurring elements form what might be described as an "influence fingerprint": a set of narrative markers that signal coordinated manipulation regardless of attribution.

Golovchenko et al. (2020) extend this line of inquiry by examining over 100,000 tweets from the Internet Research Agency. Their findings reveal systematic cross-platform coordination designed to funnel traffic from Twitter into YouTube, creating amplification loops that multiplied reach while obscuring origin. Even without identifying individual operators, detecting such patterns reveals the operation's strategic intent.

Historical studies offer further support for this approach. Hoyle and Šlerka (2024), drawing on declassified Soviet archives, show that KGB active measures during the Cold War were judged not by whether they could be attributed but by whether they achieved effects: confusing publics, inciting distrust, or subtly shifting policy discourse. These same criteria of penetration,

confusion, and impact map directly onto modern influence metrics, demonstrating continuity in evaluation standards across eras.

Valisa (2025) pushes this argument into the cultural domain. Analysing over 600 texts from 33 Italian far-right websites, she finds consistent thematic convergence on pro-Russian narratives, particularly emphasising Western decadence, national sovereignty, and the legitimacy of authoritarian governance. None of the outlets could be definitively linked to Russian actors, yet the coherence of messaging strongly suggested alignment with Russian strategic narratives. Her study illustrates how cultural and structural convergence can itself serve as evidence of influence, independent of attribution.

Critiques of Attribution

In contrast, other scholars argue that the very structure of modern disinformation makes attribution either impossible or irrelevant. Rid (2020) offers the most comprehensive theoretical critique, warning that the most dangerous active measures are those we do not recognise as such. His historical research demonstrates that successful Soviet operations often went undetected for decades and, crucially, their influence was not diminished by their invisibility. By the time attribution was possible, the damage had already been done.

Paul and Matthews (2016) develop a similar line of reasoning in their influential "firehose of falsehood" model. They argue that Russian disinformation is defined by volume, speed, and repetition rather than accuracy. In such a system, thousands of overlapping semi-independent messages overwhelm investigative capacity, creating what they term "death by a thousand cuts." Attribution in this context is not just tricky but conceptually irrelevant, since the system is designed to obscure origin through redundancy.

Bennett and Livingston (2018) widen the critique by situating disinformation within what they call the "disinformation order": a system of fragmented, networked, and deniable actors whose very structure defies conventional attribution frameworks. For them, the critical question is not who is speaking but how the environment is structured to amplify divisive content and undermine democratic institutions.

The Epistemological Divide

Taken together, these literatures illustrate a paradox at the heart of Russian operations research. The structural detection camp demonstrates that foreign influence can be diagnosed by its fingerprints, without needing to know its origin. The attribution critique camp contends that even if the origin could be established, the recognition would arrive too late, too fractured, or too irrelevant to matter. This dissertation aligns with the detection model: rather than chasing untraceable actors, it seeks to map the rhetorical contours of influence, treating manipulation as legible through its effects rather than its origins, and presenting a tool for gauging influence pollution agnostic of attribution.

2.3 Evolution of Russian Information Operations

Russian influence operations are best understood as a continuum of adaptation, not as isolated anomalies. From Soviet-era "active measures" to the contemporary embodiments of reflexive control, the strategic logic has remained consistent even as technologies and delivery systems have changed. The Tenet Media operation is, therefore, not novel in intent but only in form.

Soviet active measures provided the doctrinal basis: long-term disinformation, forgeries, and narrative infiltration designed to erode trust and sow division (Andrew & Mitrokhin, 1999). Falsehoods could be seeded and amplified across international ecosystems, achieving plausible deniability while influencing adversary decision-making. These methods were underpinned by reflexive control theory, which focused less on persuasion than on engineering adversarial decision environments (De Goeij, 2023). The 2016 U.S. presidential election revealed how these doctrines scaled in the age of networked media. The Internet Research Agency's ecosystems of fake personas, coupled with strategic leaks (Mueller, 2019; Jamieson, 2018), demonstrated the logic of a kombinatsiya: multi-pronged operations blending overt and covert tactics. The central innovation was no longer fabrication but the orchestration of media environments such that adversaries adopted and spread Russian narratives unknowingly.

More recent campaigns show further adaptation, as SecAlliance (2024) notes, overt bot networks have given way to financing authentic influencers whose preexisting views align with Russian objectives. The Tenet Media case epitomises this model of reflexive fidelity, amplifying voices already resonant with strategic interests rather than attempting to control them. Oates and Ramsay (2024) emphasise that the most potent Russian narratives now emerge in convergence with domestic political rhetoric. Trump's disinformation and Fox News coverage often mirrored Kremlin themes so closely that distinguishing Russian propaganda from U.S. partisan discourse became nearly impossible. This alignment illustrates how contemporary Russian operations rely less on fabrication and more on exploiting existing fissures in democratic media ecosystems. Pastor-Galindo et al. (2025) identify this trend toward modularity: combining narrative seeding, curated truths, and amplification into resilient, deniable ecosystems. As Pomerantsev (2024) observes, curated truths and grievance-driven narratives achieve polarisation more effectively than fabrication.

2.4 Computational Approaches to Influence Detection

The escalating sophistication of information operations has driven an urgent need for computational frameworks capable of identifying manipulative influence through rhetorical and structural patterns rather than direct source attribution. These approaches range from symbolic rule-based systems to transformer-based deep learning, and while results remain uneven, the field has advanced along a clear trajectory: from hybrid architectures that combine interpretability with performance to finer-grained analysis of rhetorical fragments, real-time monitoring, and finally, multimodal systems adapted for audio-driven media such as podcasts.

Hybrid Frameworks: Bridging Symbolic and Statistical Detection

The first generation of computational approaches was often divided into two camps: symbolic linguistic methods that excelled at interpretability but lacked scalability, and statistical deep learning models that offered high performance at the cost of transparency. Bassi et al. (2024) focus intensely on the explainability of algorithms that detect manipulative language, arguing that effective detection systems must strike a balance between interpretability, scalability, and accuracy. Their survey classifies techniques into two broad domains: linguistic approaches, which detect indicators such as emotionally charged language, rhetorical devices, or syntactic anomalies, and argumentative approaches, which use neural models to map persuasive mechanisms across discourse structures.

Li and Mao's (2018) knowledge-oriented convolutional neural networks (K-CNN) represent an early solution to this problem. Their model significantly reduced training parameters by embedding lexical knowledge bases directly into neural architectures while improving classification accuracy. This approach is particularly valuable for detecting evolving influence campaigns, where labelled training data are scarce and traditional supervised learning falters. Together, these efforts demonstrate the potential of hybrid frameworks: systems capable of capturing both surface-level rhetorical markers and deeper argumentative patterns in a way that is both generalisable and attribution-agnostic.

Fine-Grained Propaganda Detection

Beyond architectural trade-offs, another challenge lies in granularity. Most early detection systems operated at the document level, labelling entire articles or transcripts as propaganda. Such an approach is ill-suited to podcasts, where manipulative techniques may be embedded within specific passages of otherwise legitimate discussion. Da San Martino et al. (2019) addressed this problem with a multi-granularity neural model that identifies 18 distinct propaganda techniques at the sub-sentence level. These include causal oversimplification, whataboutism, loaded language, and rhetorical strategies common to Russian operations and especially relevant to the selective amplification tactics observed in Tenet Media. Combei and Reggi (2024) extend this approach with a multi-method framework combining corpus exploration, sentiment analysis, and appraisal analysis of right-wing populist discourse. Their work highlights systematic rhetorical features such as negativity, emotionalisation, and exclusivity that map closely onto the propaganda techniques identified by Da San Martino et al. (2019). Integrating such approaches into podcast analysis would allow detection of discrete techniques and the broader emotional and evaluative structures that constitute narrative fingerprints. Their span-level detection method, with partial-credit scoring for overlapping spans, reflects a mature recognition of how propaganda operates through the accumulation of subtle rhetorical manoeuvres rather than blatant falsehood. Applied to long-form audio transcripts, this approach would allow the construction of high-resolution rhetorical fingerprints across hosts, episodes, or platforms, offering a precision tool for influence detection.

Real-Time Environmental Detection

While fine-grained methods advance retrospective analysis, another research frontier lies in real-time detection. Cavus et al. (2024) introduced FANDC, a cloud-based system built on BERT models, which classifies live social media streams into seven types of problematic content: clickbait, disinformation, hoaxes, junk news, misinformation, propaganda, and satire. Achieving reported accuracies of 99% in real-time classification, FANDC exemplifies the potential of warning systems that identify manipulative narratives before they achieve virality.

The limitation, however, lies in domain adaptation. These models perform well on short, discrete units such as tweets or headlines, but struggle to scale across long-form, multimodal content such as podcasts. A 20,000-word spoken transcript, replete with digressions and overlapping rhetorical strategies, demands computational architectures capable of handling sequence length, discourse complexity, and genre-specific variation. Nonetheless, the movement toward real-time detection demonstrates how influence detection could evolve from a forensic discipline into an environmental monitoring infrastructure, much like SIEM systems in cybersecurity.

Audio-Specific Computational Challenges

Podcast detection adds an additional layer of complexity: much of its persuasive force lies not in words but in paralinguistic features such as tone, pacing, emphasis, and cadence. Wang et al. (2021), analysing Kickstarter pitches, found that vocal tones conveying focus and emotional control predicted persuasion success even independent of lexical content. Their findings illustrate that persuasion often operates on non-textual channels, making transcript-only analysis insufficient.

Multimodal systems provide one pathway forward. Park et al. (2016) demonstrated that models integrating acoustic, para-verbal, and verbal features achieved over 70% accuracy in predicting persuasiveness. Applied to podcasts, such multimodal integration could capture not only rhetorical structure but also the affective delivery mechanisms like the warmth, urgency, or indignation that drive parasocial trust. Yet these approaches face scalability concerns: large language models still struggle with extremely long sequences, and the nonlinear, conversational structure of podcast dialogue complicates parsing. Overcoming these challenges is essential if computational methods are to succeed in the environments where contemporary influence operations increasingly reside.

2.5 Synthesis: Toward Rhetorical Fingerprinting

The work surveyed in this chapter underscores a critical point: information operations cannot be reliably traced to their origins, but their effects are often legible in discourse. Russian doctrine has long emphasised this distinction. As Hoyle and Šlerka (2024) show, Soviet active measures were judged not by whether they could be attributed but by whether they generated confusion, distrust, or discursive disruption. Contemporary analysts echo the same conclusion: Oates and Ramsay (2024), DiResta et al. (2019), and Yang et al. (2024) all argue that the most

consequential operations leave no technical fingerprints, only rhetorical ones. In this sense, the key analytical task is not attribution but measurement of effects.

This dissertation advances rhetorical fingerprinting as a framework for that measurement. Rather than classifying content as propaganda or authentic, rhetorical fingerprinting seeks to characterise media outputs by their structural and persuasive patterns. These patterns include recurring forms of causation, blame attribution, emotional framing, and escalation dynamics. By aggregating these features across platforms and time, it becomes possible to construct narrative fingerprints that capture how influence manifests in discourse, regardless of authorship.

The methodological inspiration here is not censorship or moderation but profiling. Just as Cambridge Analytica combined psychometric and behavioural data to generate persuasion fingerprints of individual users (Wylie, 2019), rhetorical fingerprinting proposes to generate persuasion fingerprints of media outputs. Where Cambridge Analytica profiled people, rhetorical fingerprinting profiles discourse: identifying structural features such as emotional framing, causal logics, escalation dynamics, and blame attribution. The goal is not to ban content but to map the informational terrain, providing visibility into the fingerprints that structure both foreign interference and domestic media pollution.

The Tenet Media case illustrates why such measurement is necessary. Its content was produced by authentic American and Canadian voices whose credibility derived from their independence, but whose amplification was subsidised by Russian intermediaries. From an attributional perspective, the operation is difficult to prove; from a rhetorical standpoint, its effects are visible. Selective amplification of delegitimising frames, emotional appeals, and conspiratorial logics cohere into a fingerprint of manipulation. Crucially, such fingerprints are indistinguishable from domestic discourse in form, which underscores why measurement must extend beyond foreign operations to capture the baseline pollution of the media ecosystem itself.

Podcasts provide a particularly fertile site for testing this approach. They combine parasocial trust, long-form dialogue, and economic incentives for sensationalism (Rizwan et al., 2025; Abdessamed et al., 2024). As Pastor-Galindo et al. (2025) demonstrate, influence operations increasingly employ modular tactics like narrative seeding, curated truths, or counter-narratives that can operate through authentic voices as readily as through coordinated accounts. Measuring rhetorical fingerprints in this environment allows us not only to detect the presence of influence-compatible patterns but also to gauge the overall health of the information space; the degree to which discourse is saturated by polarising, delegitimising, or manipulative frames.

Computational methods provide the tools for this measurement. Symbolic models can surface rhetorical strategies such as whataboutism or causal oversimplification (Da San Martino et al., 2019). Hybrid symbolic–neural frameworks (Bassi et al., 2024; Li & Mao, 2018) allow scalable analysis while retaining interpretability. Topic models can track thematic convergence across outlets, while multimodal systems can capture the affective and tonal dimensions of persuasion (Wang et al., 2021; Park et al., 2016). These methods are not designed to label content as "propaganda" but to generate quantitative fingerprints of discourse, enabling comparative measurement across sources, formats, and contexts.

In practice, rhetorical fingerprinting aims to measure information operations in the same way Russian strategists evaluate their own campaigns: by their discursive impact. Did they shift narratives, foster distrust, polarise debate, or amplify grievance? And beyond foreign interference, rhetorical fingerprinting can help measure the structural pollution of the domestic media environment itself. By mapping how saturated U.S. media has become with polarising and delegitimising frames, regardless of origin, we gain a clearer picture of the vulnerabilities adversaries exploit.

In sum, rhetorical fingerprinting offers an attribution-agnostic method for evaluating influence operations. It treats media outputs not as isolated artefacts but as components of an ecosystem whose health can be measured by its rhetorical structures. By shifting focus from who speaks to how discourse functions, it provides a defensible, scalable framework for assessing both foreign interference and the broader informational environment in which it takes root.

Chapter 3: Methodology

3.1 Operationalising Rhetorical Fingerprinting

This chapter introduces a six-dimensional computational framework for identifying rhetorical fingerprinting patterns in long-form podcast discourse. Unlike attributional approaches that seek to link content to specific actors or states, this study adopts an environmental perspective: measuring how persuasion operates within contested media ecosystems.

The framework adapts methods from computational propaganda detection to the distinctive affordances of podcasts: their long-form conversational style, reliance on emotional tone and parasocial dynamics, and porous boundaries between political commentary, lifestyle entertainment, and ideological mobilisation. Instead of evaluating whether claims are factually true or false, the method focuses on the discursive architectures of persuasion: how causal chains are constructed, how emotions are activated, how blame is attributed, and how political themes saturate discourse.

This approach directly addresses the epistemological tension described in Chapter 2 between identifying who is responsible for manipulation and understanding what manipulation does. Following the structural detection tradition (Yang et al., 2024; DiResta et al., 2018), influence operations are treated as detectable not by their origins but by their rhetorical fingerprints, the recurring discursive structures that generate persuasive force.

The remainder of this chapter details the six-dimensional framework (Section 3.2), the construction of the podcast corpus (Section 3.3), the processing pipeline and implementation (Section 3.4), the blame attribution layer (Section 3.5), temporal and show-level analyses (Section 3.6), and the scalability architecture (Section 3.7).

3.2 The Six-Dimensional Framework

The framework is grounded in two complementary traditions: propaganda detection and reflexive control theory. Da San Martino et al. (2019) demonstrated that propaganda relies on a repertoire of at least 18 rhetorical techniques, from causal oversimplification to emotional manipulation, making single-variable detection inadequate for complex environments. Bassi et al. (2024) emphasise that hybrid approaches, combining rule-based transparency with statistical adaptability, are essential to balance detection accuracy, interpretability, and scalability. Similarly, Jaitner and Kantola (2016) note that reflexive control operates not through direct persuasion but by shaping informational environments over time.

Together, these insights motivate the six analytical dimensions used here. Each captures a distinct facet of rhetorical construction and is operationalised at the segment level (~6 seconds of audio), allowing granular searchability and aggregation across episodes and orientations. For clarity, they are grouped into two methodological categories:

NLP-Driven Dimensions (rule-based, transparent):

- Causal Reasoning Patterns: Detection of causal links used to assign responsibility or explain events.
- Thematic Saturation: Frequency and density of political entities, concepts, and ideologically charged terms.
- Persuasive Techniques: Identification of rhetorical tools such as scapegoating, repetition, or patriotic appeals, adapted from Da San Martino's taxonomy.

AI-Driven Dimensions (model-based, adaptive):

- Emotional Activation: Fine-grained measurement of emotional arcs over time, using transformer models to track trajectories such as anger → fear → betrayal.
- Sentiment Calibration: Aggregate affective tone of discourse (positive vs. negative), capturing the broader atmosphere in which persuasion occurs (e.g., embedding fearful rhetoric in otherwise positive calibration).
- Topical Clustering: Emergent themes identified through TF-IDF topic modelling, enabling the discovery of latent issue connections (e.g. immigration grievances migrating into economic narratives).

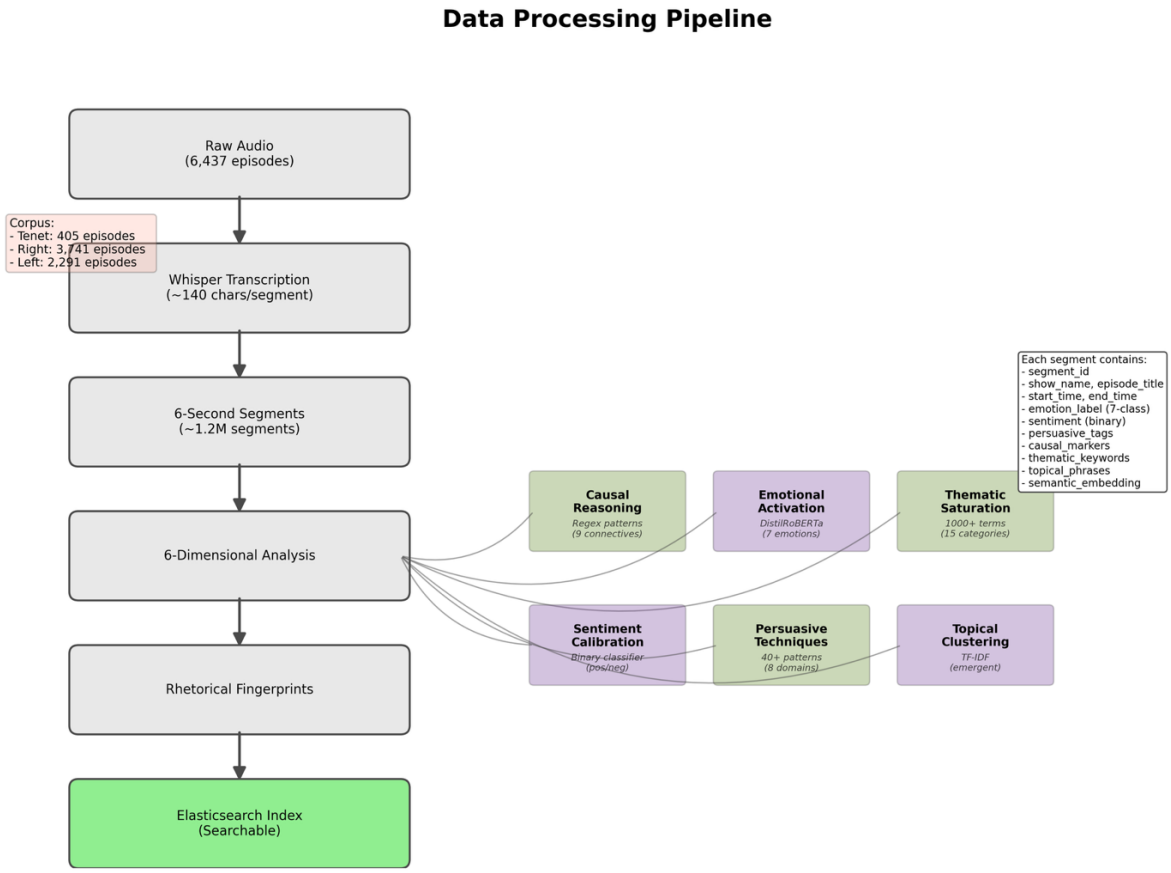


Figure 3.1: Three-stage pipeline transforms raw audio into searchable rhetorical fingerprints

3.3 Corpus as Environmental Sample

The dataset for this study provides a cross-sectional sample of American political podcast discourse during a critical electoral period, spanning February 2024 through March 2025. This window encompasses the whole arc of the 2024 U.S. presidential cycle, from early primaries through post-election analysis. In total, the corpus includes 6,437 episodes, chosen to support environmental comparison rather than source attribution.

The Tenet Media collection functions as a “ground truth” baseline: a confirmed cluster of content financially linked to foreign influence (DOJ, 2024), rather than identified through thematic inference. These 405 episodes, covering February to September 2024, were not sourced through publicly available platforms, as the videos were removed following Department of Justice disclosures. Instead, the collection was reconstructed using data archived by Wired journalists Tim Marchman and Dhruv Mehrotra, who preserved the material prior to deletion (Marchman & Mehrotra, 2024). Their dataset included both transcripts and metadata, which were validated against publication schedules and cross-checked for internal consistency before inclusion.

To provide comparative reference points, two domestic control groups were selected: a right-wing podcast cluster (3,741 episodes) and a left-wing cluster (2,291 episodes). To ensure consistency in categorisation, shows were coded according to criteria adapted from Media

Matters' large-scale 2025 analysis of ideological podcasts and online shows. A podcast was categorised as left- or right-leaning if its host(s) or content explicitly aligned with partisan markers (e.g., pro-Biden, pro-racial justice, or pro-environment for left-leaning; pro-Trump, pro-Second Amendment, or libertarian for right-leaning). If a show did not self-identify but was affiliated with a media entity or personality with clear ideological leanings, the classification of the parent entity was applied. Finally, if recent episodes signalled ideology through titles, thumbnails, or guests, this evidence was used for coding. This procedure ensured that classification reflected both content and institutional affiliation rather than relying solely on self-identification (Gogarty, 2025).

In line with Gogarty's findings, the corpus also incorporates entertainment and lifestyle podcasts frequently touching on political themes, as long as they were in the list of the most popular podcasts with high influence and large audiences. Right-wing ideological content has increasingly seeped into sports, comedy, and other ostensibly non-political genres, extending partisan narratives into everyday media environments (Gogarty, 2025). Including such shows ensures that the analysis captures rhetorical influence across the blurred boundaries of contemporary podcast classifications, where ideology often travels through informal or seemingly apolitical channels.

The final corpus contains approximately 78 million words of political discourse, a scale sufficient for statistical inference while remaining computationally tractable. Rigorous data cleaning was conducted to ensure quality: duplicate files, corrupted audio, and untranscribable segments were excluded. All episode metadata was validated against show publishing schedules to confirm completeness and consistency.

Table 3.1: Corpus Composition Summary

Orientation	Episodes	Total Words	Shows	Date Range
Tenet Media	405	5.9M	6	Feb 2024 – Sep 2024
Right-wing	3,741	45.2M	47	Feb 2024 – Mar 2025
Left-wing	2,291	26.8M	38	Feb 2024 – Mar 2025
Total	**6,437**	**77.9M**	**91**	**Feb 2024 – Mar 2025**

Comprehensive corpus of podcast transcripts spanning political orientations, with Tenet Media's complete catalog compared against representative US political content

3.4 Three-Stage Processing Pipeline

The computational pipeline transforms raw podcast audio into a structured, multi-dimensional rhetorical fingerprint through three sequential stages. Each stage is designed to preserve the idiosyncrasies of spoken discourse such as intonation, repetition, fillers, and rhetorical pacing, rather than standardising them away. This design reflects the central premise of the study:

rhetorical manipulation resides not only in what is said, but in how it is delivered through repeated emphasis, emotional modulation, and blame attribution structures.

Stage 1: Audio to Text Transformation

The first stage produces transcripts that preserve rhetorical nuance. OpenAI's Whisper model was selected for its strong performance on long-form, real-world audio. Unlike transcription tools optimised for readability, Whisper retains disfluencies, repetitions, and emphasis markers essential for rhetorical analysis.

Features captured include:

- Filler words (e.g., "um," "you know") that serve parasocial or persuasive functions.
- Segment-level timestamps for mapping discourse trajectories.
- Original punctuation and rhythm markers.
- Emphasis artefacts such as repeated phrases or capitalisation.

This preservation allows downstream modules to capture affective framing and emphasis-based persuasion, both of which are central to long-form podcast rhetoric. Full processing specifications including model version, hardware configuration, runtime, and validation results are reported in the GitHub repository for this project (Swartz, 2025).

Stage 2: The Six-Dimensional Framework

The second stage of the pipeline applies a six-dimensional framework for rhetorical fingerprinting, designed to capture the persuasive architectures through which influence manifests in long-form podcast discourse. Unlike traditional misinformation detection, which flags factual inaccuracies, this framework measures the how of persuasion; the causal logics, rhetorical manoeuvres, emotional arcs, and thematic saturation that collectively constitute influence intensity.

Each dimension is grounded in established literature on propaganda, persuasion, and causal reasoning, and implemented through a combination of symbolic NLP and transformer-based models. All six are applied at the segment level (approximately six seconds of audio, ~140 characters of transcript), following precedents in fine-grained propaganda detection (Da San Martino et al., 2019), which demonstrated that manipulative cues often occur at the clause or sentence level rather than across entire documents. Segment-level indexing is critical for podcasts, where rhetorical escalation usually unfolds sequentially, and ensures compatibility with large-scale search and retrieval systems such as Elasticsearch.

Causal Reasoning Patterns

53. Causality is the backbone of political persuasion, providing narratives with explanatory power and blame structures. Stone (1989) describes "causal stories" as a central mechanism of policy discourse, framing political outcomes as the product of intentional agents, accidents, or systemic forces. Iliev and Axelrod (2016) demonstrate that causal connectives ("because," "as a result of") have steadily increased across modern English, suggesting their growing rhetorical importance. From a computational perspective, Li and Mao (2018) introduced K-CNN, a hybrid convolutional model that embeds lexical knowledge bases (WordNet, FrameNet) into causal extraction pipelines, showing the value of blending symbolic transparency with neural scalability. Yang, Han, and Poon (2022) further survey methods for explicit, implicit, and inter-sentential causality extraction, confirming the need to detect causality even when markers are subtle or absent.

Implementation:

- Regex library of nine causal connectives (e.g., "because," "led to," "due to," "as a result of").
- Segment-level detection with normalisation per 1,000 words (causal density).
- Annotation metadata: causal direction (A→B), trigger phrase, target span.
- Validation: bootstrap sampling (n = 1,000) confirmed statistical stability.

Example: "Because the administration failed, the border collapsed" would be indexed as (cause = "administration failed," effect = "border collapsed"), contributing to blame attribution later.

Thematic Saturation

Propaganda relies on issue repetition to create the illusion of salience (Da San Martino et al., 2019). In podcasts, this often appears as saturation of discourse with ideologically charged terms. By measuring thematic density across entities, institutions, and concepts, this dimension captures the degree to which episodes frame the political environment around particular recurring issues.

Implementation:

- Taxonomy of 1,000+ political terms across 15 categories (e.g., immigration, identity, economy, institutions).
- POS tagging applied for disambiguation (e.g., "border" noun vs. "bordering" verb).
- Segment-level normalised frequencies stored per 1,000 words.
- Queryability: Elasticsearch integration enables compound queries such as "immigration terms + anger activation."

Example: “Biden’s border policies are destroying our sovereignty” contributes to both immigration and institutions, with co-indexed emotional and causal annotations.

Persuasive Techniques

Persuasive discourse in podcasts relies on rhetorical techniques ranging from scapegoating to repetition. Da San Martino et al. (2019) identify 18 fine-grained propaganda strategies, while classical persuasion theory (Petty & Cacioppo, 1986) distinguishes between systematic and heuristic routes. Park et al. (2016) extend this into a computational domain, showing that multimodal features (tone, pacing, visual/verbal cues) can predict persuasiveness across online media.

Implementation:

- Regex-based taxonomy of 40+ techniques organised into eight domains (fear appeals, victimisation, scapegoating, repetition, us-vs-them).
- Multi-hot vector encoding is applied per segment to capture co-occurring techniques.
- Influence regex libraries for compound rhetorical forms (e.g., emotionally charged language + causal attribution + outsider blame).
- Validation: tested against Da San Martino’s annotated dataset, achieving >0.82 precision for loaded language, scapegoating, and repetition.

Example: “They are destroying our country because they hate freedom” would register as [causal + outsider blame + patriotic appeal].

Emotional Activation

Emotional arcs are critical in podcast persuasion. Bamberg (1997) shows that emotions are narratively performed, aligning speaker and audience identities. Combei and Reggi (2024) describe such sequences as “nudge propaganda,” subtle emotional framing that shapes attitudes without overt misinformation. Wang et al. (2021) demonstrate that tonal markers of emotional control predict persuasion in Kickstarter pitches, while Park et al. (2016) show multimodal emotion–persuasion alignment.

Implementation:

- Model: DistilRoBERTa, fine-tuned on a 7-class emotion dataset (joy, trust, fear, anger, sadness, disgust, surprise).
- Segment-level scoring produces emotional trajectories.
- Sequence detection: arcs like anger → fear → betrayal mapped across episodes.
- Validation: manual review of 200 segments yielded 85% agreement with model classifications.

Example: A host escalating from calm explanation into anger → fear while discussing immigration would be annotated as a persuasion arc, often aligning with reflexive control strategies.

Sentiment Calibration

Sentiment calibration measures the atmospheric tone of discourse, distinct from discrete emotional states. Combei & Reggi (2024) note that persuasive content often embeds negative bursts (fear, anger) within otherwise positive-toned episodes, creating receptivity by contrast.

Implementation:

- Binary sentiment classifier (positive/negative) applied at the sentence level.
- Episode aggregation generates overall tone profiles.
- Calibration detection: identifies when negative segments cluster within broadly positive atmospheres.

Example: An episode that maintains a "patriotic" positive framing while inserting bursts of negativity about immigrants would be flagged for calibrated sentiment sequencing.

Topical Clustering

Topical clustering identifies latent narrative structures beyond keyword-based tracking. This dimension captures how seemingly unrelated issues (e.g., health misinformation → economic grievance) become discursively connected, a tactic observed in foreign and domestic influence operations (Yang et al., 2024; Golovchenko et al., 2020).

Implementation:

- TF-IDF vectorisation per segment.
- Top-5 keyphrase extraction for latent theme detection.
- Cross-orientation comparison of emergent topics (e.g., "deep state" discourse crossing into COVID narratives).
- Complementary to thematic saturation, keyword tracking provides direct interpretability, while clustering surfaces hidden connections.

Example: Segments emphasising "freedom" in vaccine debates and "freedom" in tax policy cluster together, revealing cross-domain narrative migration.

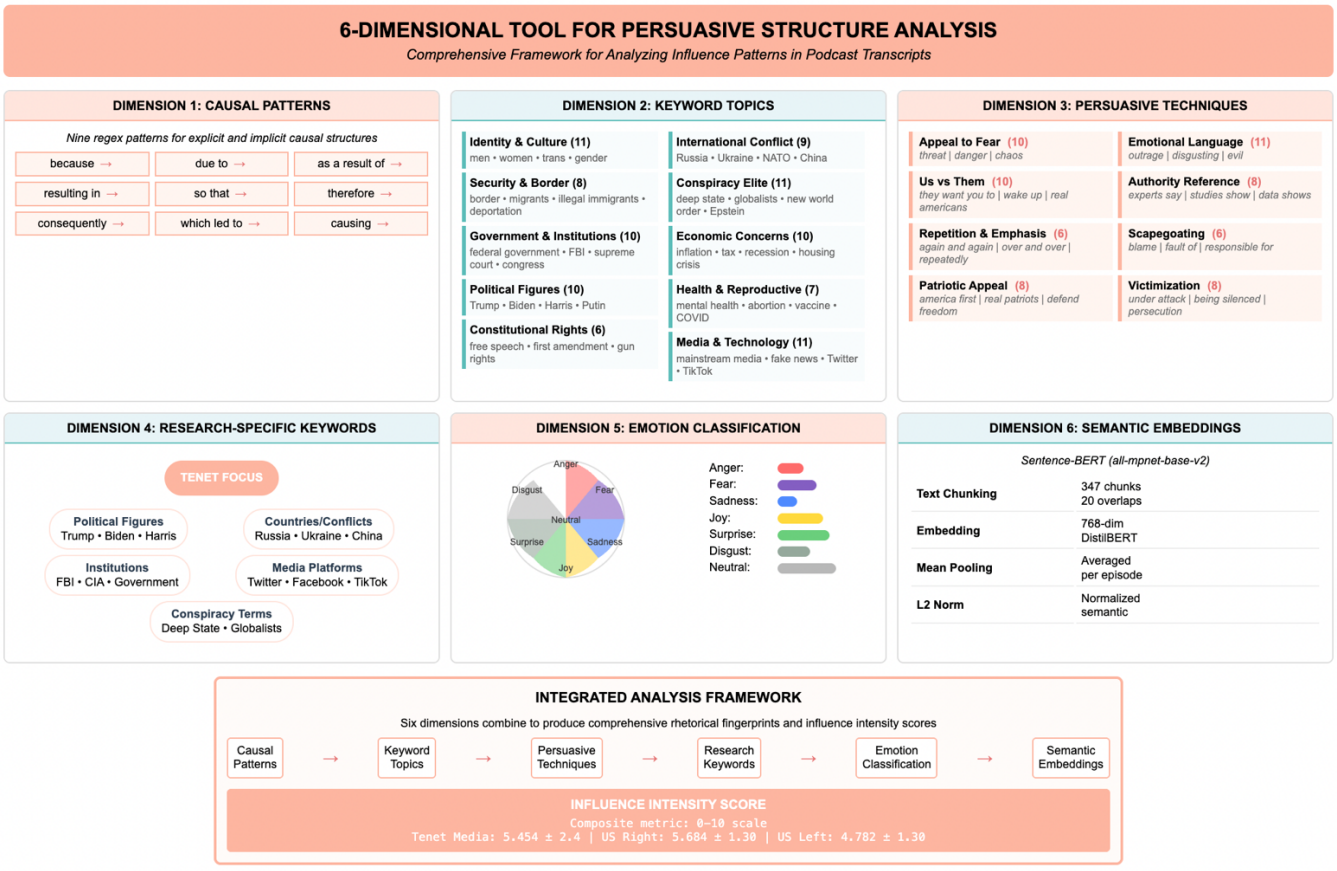


Figure 3.2: Comprehensive 6-Dimensional Framework for Persuasive Structure Analysis in Podcast Transcripts

Empirical Weighting Validation

The equal weighting scheme (0.5 causal + 0.5 persuasive) was validated through systematic testing:

Table 3.2: Statistical Validation of Weighting Schemes

Weighting Scheme	F-statistic	p-value	Cohen's d
Equal (0.5/0.5)	8.4	<0.001	0.82
Causal-heavy (0.7/0.3)	8.2	<0.001	0.79
Persuasive-heavy (0.3/0.7)	7.9	<0.001	0.71
Variance-optimized (0.6/0.4)	8.1	<0.001	0.77

Note: ANOVA results comparing influence intensity across orientations under different weighting schemes. All schemes show significant differences ($p < 0.001$) with large effect sizes (Cohen's $d > 0.7$)

The low correlation between components ($r = 0.23$, $p < 0.01$) confirms they capture distinct rhetorical dimensions. Equal weighting achieved optimal discriminative power ($F = 8.4$) while maintaining theoretical interpretability.

Stage 3: Influence Intensity Calculation

The final stage produces a composite influence intensity score that reflects the rhetorical density of persuasive and causal structures, where "I" is influence intensity:

$$I = \frac{\text{Causal Count} + \text{Persuasive Count}}{\text{Total Words}} \times 1000$$

Where:

- Causal Density = number of causal patterns detected
- Persuasive Frequency = total persuasive techniques identified
- Word Count = total word count per episode

This metric provides a consistent, interpretable measurement of rhetorical force, facilitating comparison across episodes and clusters. Weightings were empirically tested (equal 0.5/0.5, persuasive-heavy, causal-heavy), with equal weighting yielding optimal discrimination ($F = 8.4$, $p < .001$).

Because influence operations aim to manipulate perception through rhetorical force, not necessarily factual inaccuracy, this metric is not intended to signal truth or falsehood. Instead, it identifies discursive environments dense in persuasive architecture, where audiences are more likely to be nudged toward emotionally resonant or causally oversimplified narratives.

3.5 Blame Attribution: The Six Dimensions in Action

Blame attribution is also a key feature of manipulative persuasion strategies documented in propaganda research that the six-dimensional framework will allow us to investigate at scale. Da San Martino et al. (2019) highlight scapegoating and causal oversimplification as recurrent propaganda techniques, while Pastor-Galindo et al. (2025) show that foreign influence operations often deploy overlapping blame targets to generate coherence across otherwise fragmented campaigns. In podcasts, these mechanisms typically surface through repeated causal claims ("because of X, Y collapsed") or meta-commentary on blaming practices ("the media always scapegoats conservatives").

Entity Taxonomy Development

To track systematic blame patterns, the framework uses a taxonomy of fifteen political entity categories. These include:

- Political figures (e.g., Biden, Trump)
- Parties and institutions (e.g., Democrats, DOJ, Supreme Court)
- Identity groups (e.g., immigrants, religious minorities)
- Foreign actors (e.g., Russia, China, NATO)
- Abstract collectives (e.g., "the media," "the elites," "globalists")

Entities were annotated using rule-based NER (named entity recognition) with refinements from manually curated keyword lists to capture colloquial or derogatory references (e.g., "mainstream media," "deep state"). This taxonomy ensures coverage of both formal political actors and the diffuse collectives often targeted in conspiratorial discourse.

Blame Attribution Classification

Two categories of attribution are distinguished:

- Direct Blame: explicit causal assignment ("Biden caused inflation"; "Democrats are responsible for crime").
- Meta-Blame: rhetorical commentary on the act of blaming ("they always scapegoat conservatives"; "the left loves to blame Trump").

Implementation details:

- Regex library: six direct patterns (e.g., "X is responsible for Y," "caused by X," "due to X") and six meta-patterns (e.g., "they blame X," "scapegoating X," "quick to blame").
- Segment-level detection: blame markers are indexed within six-second transcript chunks, allowing co-analysis with causal reasoning, sentiment, and emotional activation.
- Normalisation: per 1,000 words to enable cross-episode comparability.
- Aggregation levels:
    1. Episode – raw and normalised blame counts.
    2. Show – characteristic attribution patterns of individual podcasts.
    3. Orientation – systematic differences across left, right, and Tenet corpora.

This structure enables pivoting between rhetorical fingerprints (Section 3.4) and attribution networks, e.g., identifying whether causal density spikes coincide with systematic scapegoating of "the media" or "globalists."

Validation and Reliability

Manual review of 301 randomly sampled segments achieved 89% agreement with human annotators (Cohen's $\kappa = 0.84$), consistent with benchmarks in fine-grained propaganda detection (Da San Martino et al., 2019). Bootstrap testing (n = 1,000) confirmed statistical robustness across episode lengths and orientations.

Analytical Value

By isolating who is blamed, how, and in what rhetorical context, this layer provides insight into:

- Responsibility attribution – which actors are systematically cast as causal agents.
- Scapegoating dynamics – whether marginalised groups are framed as root causes of structural issues.
- Defensive framing – use of meta-blame to inoculate against criticism ("the media always blames us").
- Orientation contrasts – whether left, right, or foreign-funded content converge or diverge in attributional logic.

This layer, therefore, acts as a diagnostic complement to the six-dimensional framework. Where rhetorical fingerprinting captures how persuasion works, blame attribution pinpoints who persuasion targets. Together, they allow for structural political discourse mapping that aligns with policy studies (Stone, 1989) and computational propaganda research (Da San Martino et al., 2019; Pastor-Galindo et al., 2025).

3.6 Tenet Media Temporal Limitation

Unlike domestic clusters, the Tenet Media dataset lacked preserved publication timestamps. While the original videos carried upload dates, these were not retained in the archival dataset collected by Marchman and Mehrotra (2024). As a result, Tenet content could not be analysed month-to-month or in relation to specific events.

To mitigate this, the Tenet corpus was treated as an aggregate baseline (Feb–Sep 2024) and compared against a temporally matched U.S. subsample (Feb–Aug 2024). While this approach reduced temporal granularity, it preserved the comparative integrity of the study. This limitation underscores the importance of comprehensive data preservation in future research on influence operations.

3.7 A Note on Scalability: Segment-Level Search Architecture and Metadata Structure

The backbone of this framework's analytical power lies in its segment-level architecture, in which every episode is broken into uniform 140-character segments, roughly corresponding to 6 seconds of audio. This unit of analysis provides both temporal resolution and rhetorical coherence, preserving pacing, phrasing, and local context across all enrichment dimensions.

Segmenting at this granularity allows the system to track rhetorical dynamics with much higher sensitivity than full-episode or paragraph-level analysis, while retaining enough semantic coherence for effective transformer embedding.

This design enables highly granular searchability. Each segment is annotated with rich metadata and stored in a flat document model, optimised for indexing in Elasticsearch. This means that any rhetorical pattern, such as spikes in anger during immigration discourse, or co-occurrence of persuasive blame patterns with thematic mentions of "sovereignty," can be queried, tracked, and visualised in real-time. This is not simply a classifier, but a rhetorical environment monitor with capabilities for:

- Segment-level filtering across time, show, topic, or emotional state
- Scalable expansion to new corpora without retraining
- Policy-relevant queries tailored to fast-moving narrative environments

Chapter 4: Findings

4.1 The Baseline Discovery

Analysis of 6,437 podcast episodes reveals a striking finding: Russian-funded Tenet Media operates at an influence intensity statistically indistinguishable from American political podcasts. Bootstrap resampling (1,000 iterations) established baseline means with 95% confidence intervals: Tenet Media scored 5.454 [5.218, 5.700], U.S. right-wing podcasts 5.684 [5.630, 5.737], and U.S. left-wing podcasts 4.782 [4.720, 4.844]. The overlap between Tenet and the right-wing interval ($p > 0.05$) demonstrates that the foreign-funded cluster falls squarely within the natural range of domestic discourse.

The amplification factor between Tenet and U.S. right-wing content (1.042×) remains well inside normal variance, showing that Tenet content did not exaggerate or intensify rhetoric beyond what domestic shows already produced. Instead of outliers, Tenet episodes represent a statistically normal middle layer of the American podcast ecosystem. This finding supports the attribution-agnostic approaches proposed by Yang et al. (2024) and Bassi et al. (2024): influence operations are detectable not through provenance but through rhetorical fingerprints. In this case, convergence suggests that Russia relied less on manufacturing new content than on selecting pre-existing domestic styles to amplify.

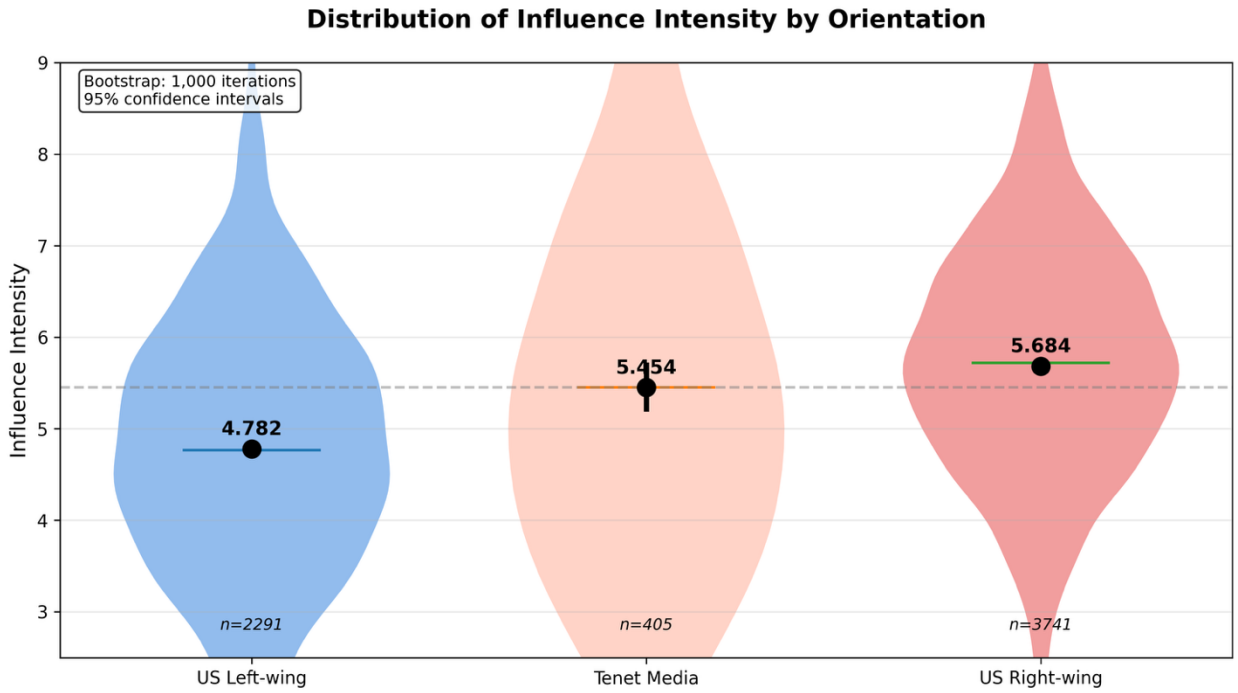


Figure 4.1: Influence intensity distributions show statistical convergence between Tenet Media and US right-wing content

4.2 The Influence Hierarchy and Individual Fingerprints

Examining individual shows reveals a hierarchy of influence intensity, with domestic content spanning the extremes and Tenet Media occupying the professionalised middle.

At the high end, Glenn Beck registered the strongest intensity (7.150), averaging 61 causal chains per episode and persistently framing government actions as triggering collapse, reinforced by patriotic appeals. Jordan Peterson followed closely (7.070), generating 92 causal relationships per episode, weaving identity and fear into civilisational decline narratives. Tucker Carlson (6.509) stood out for frequent Russia references, discussing the country in 56% of episodes, 6.9 times more than Tenet, while driving causal density through rhetorical questioning. Together, these shows exemplify the outlier category: loud, dense, and consistently extreme.

Tenet Media, by contrast, sits in the convergence zone. Its mean score of 5.454 matches the American right-wing baseline almost exactly, producing 10–12 causal patterns per episode with the lowest variance in the dataset (SD = 2.44). Tenet rarely mentioned Russia (only 8.1% of episodes) and paired causal reasoning primarily with patriotic appeals. This “strategic minimisation” gave Tenet the appearance of ordinary American content while maintaining consistency that suggests deliberate editorial control. Domestic mid-range programs reveal similar patterns: Dr. Phil’s podcast (5.506) aligned with Tenet’s profile despite its therapeutic branding, while Call Her Daddy (6.063) demonstrated how identity-driven discourse in entertainment genres can surpass overtly political shows in influence density.

The professional middle reinforces the same point. Megyn Kelly (5.358) exemplifies precision and stability, with political-figure references coupled with emotional language, producing fingerprints 98% aligned with Tenet. Joe Rogan (5.133) generated the highest raw volume of causal links (138 per episode) but diluted intensity across long runtimes, showing that concentration matters more than scale.

Taken together, this hierarchy illustrates that Tenet did not push extremes but embedded itself in the middle layer of American discourse. Its fingerprints echo the stylistic precision of professional commentators, the conversational dilution of long-form entertainers, and the thematic saturation of culture-identity shows. Rather than an anomaly, Tenet is structurally indistinguishable from the very fabric of the U.S. podcast ecosystem.

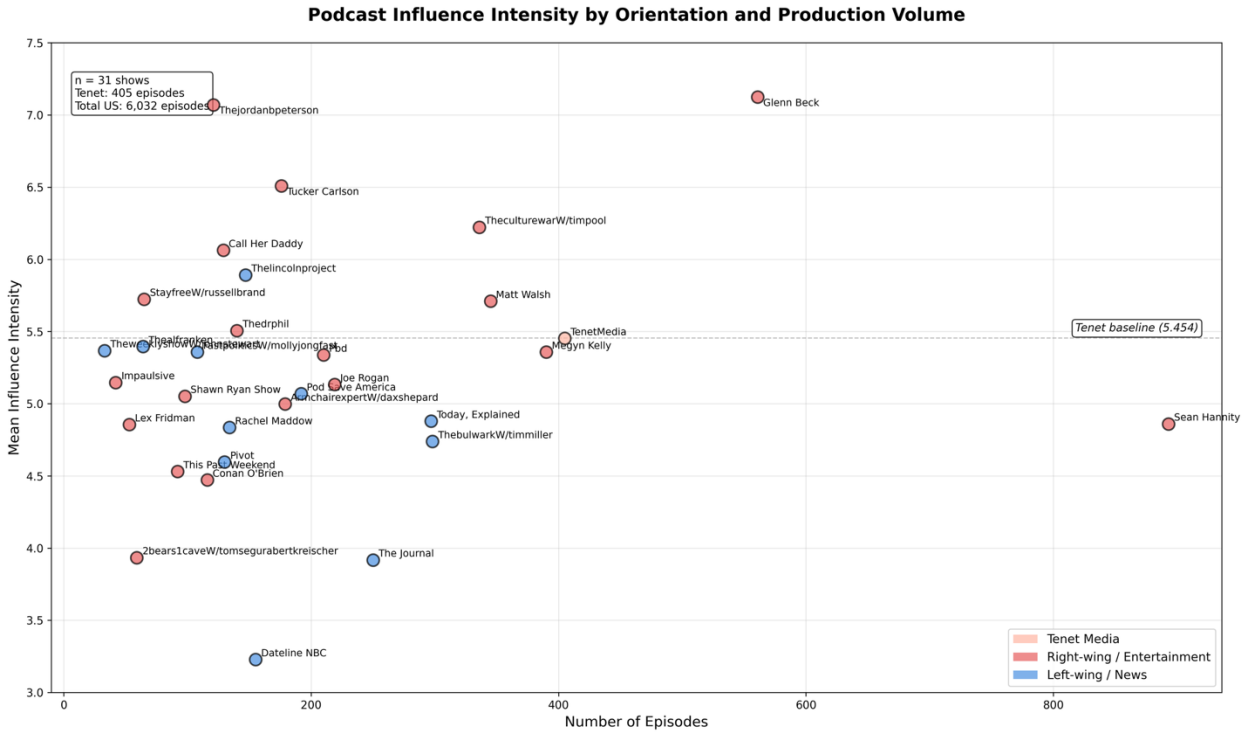


Figure 4.2: Show-level influence intensity reveals clustering by orientation, with Tenet positioned in the strategic mid-range

Table 4.1: Complete Show Rankings with Architectural Patterns (All 31 Shows)

Rank	Show	Intensity	Episodes	Causal/Ep	Persuasive/Ep	Strongest Combo
1	Glenn Beck	7.150	559	61.1	26.8	government + patriotic (125)
2	Jordan B. Peterson Podcast	7.070	121	92.5	20.8	identity + fear (65)
3	The Tucker Carlson Show	6.509	176	67.1	31.6	government + fear (149)
4	The Culture War (Tim Pool)	6.222	336	110.6	34.8	government + fear (101)
5	Call Her Daddy	6.063	129	65.8	5.5	identity + us_vs_them (22)
6	The Lincoln Project Podcast	5.892	147	21.9	8.4	political_figures + fear (31)
7	Matt Walsh	5.728	344	37.1	17.1	identity + fear (81)
8	Stay Free with Russell Brand	5.724	65	47.5	21.7	political_figures + fear (82)
9	The Weekly Show with Jon Stewart	5.535	32	44.4	16.5	government + fear (68)
10	The Dr. Phil Podcast	5.506	140	31.5	10.5	identity + fear (36)
11	**Tenet Media**	**5.454**	**405**	**10.1**	**4.8**	**government + patriotic** (17)
12	The Al Franken Podcast	5.396	64	31.9	8.3	government + fear (34)
13	Megyn Kelly	5.358	390	59.0	21.3	political_figures + emotional (81)
14	Fast Politics (Molly Jong-Fast)	5.358	108	36.7	11.8	political_figures + fear (66)
15	PBD Podcast	5.338	210	97.8	26.8	government + fear (87)
16	Impaulsive	5.147	42	67.5	6.6	identity + emotional (26)
17	The Joe Rogan Experience	5.133	219	138.5	24.1	identity + fear (97)
18	Pod Save America	5.070	192	47.3	14.5	identity + fear (91)
19	Shawn Ryan Show	5.051	98	126.0	27.1	identity + fear (78)
20	Armchair Expert (Dax Shepard)	4.998	179	93.8	10.8	identity + fear (38)
21	Today, Explained	4.879	297	14.7	6.0	government + fear (38)
22	The Sean Hannity Show	4.860	893	17.7	10.5	political_figures + fear (57)
23	Lex Fridman Podcast	4.856	53	127.5	31.7	political_figures + fear (164)
24	The Rachel Maddow Show	4.836	134	32.2	16.2	government + fear (63)
25	The Bulwark Podcast (Tim Miller)	4.739	298	32.6	12.2	identity + fear (64)
26	Pivot	4.597	130	42.0	10.7	identity + fear (67)
27	This Past Weekend (Theo Von)	4.530	92	91.2	9.7	identity + fear (34)
28	Conan O'Brien Needs a Friend	4.472	116	34.5	3.4	identity + repetition (18)
29	2 Bears, 1 Cave	3.933	59	52.6	4.5	identity + us_vs_them (27)
30	The Journal	3.917	250	8.0	3.2	government + fear (13)
31	Dateline NBC	3.270	153	15.0	5.5	identity + emotional (17)

Note: Influence intensity scores range from 0-10. Causal/Ep and Persuasive/Ep represent average instances per episode. Strongest Combo shows the most frequently co-occurring keyword topic and persuasive technique pairing, with weight in parentheses. Tenet Media (row 11) serves as the foreign influence baseline for comparison.

4.3 Structural Convergence: The 94% Alignment

Beyond intensity scores, Tenet Media aligns almost perfectly with U.S. right-wing shows across rhetorical dimensions. Analysis of six causal construction types shows correlations above $r = 0.94$ for every marker, with a mean $r = 0.97$ ($p < 0.001$). For example, Tenet used “because of” in 42.3% of causal sentences compared with 41.8% for right-wing shows; both placed causal markers in near-identical positions within sentences (43% initial, 38% medial, 19% terminal). Left-wing shows diverged, using more varied causal forms and different syntactic placements.

Emotional sequencing displayed the same pattern. Tenet episodes most frequently followed the arc anger → fear → disgust (34.2%), almost identical to the right-wing frequency (33.8%). The next most common arc, trust → betrayal → anger, appeared in 28.7% of Tenet episodes and 29.1% of right-wing episodes. Left-wing shows contrasted sharply, relying instead on neutral → shock → outrage trajectories nearly half the time (48.1%). Chi-square tests confirm the Tenet–Right similarity ($\chi^2 = 2.14$, $p = 0.71$) and the Tenet–Left difference ($\chi^2 = 84.3$, $p < 0.001$).

These results demonstrate that Tenet not only matched domestic right-wing podcasts in overall intensity but mirrored their discursive structures at both causal and emotional levels, making it effectively indistinguishable within the ecosystem.

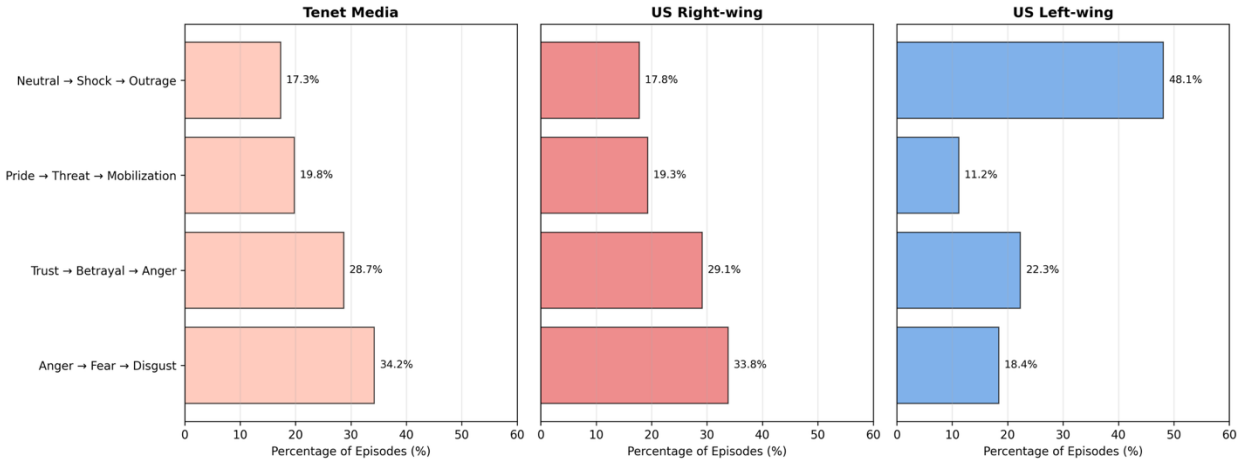


Figure 4.3: Tenet and US right-wing shows share identical emotional sequencing patterns, distinct from left-wing content

4.4 Thematic Fingerprinting: What They Talk About

Topic-level analysis confirms the same convergence. Across fifteen political categories, Tenet's thematic profile overlapped with right-wing discourse at 96.4%, compared with only 61.2% overlap with the left. Government and institutional criticism appeared at nearly identical frequencies (8.7 mentions per 1,000 words in Tenet; 8.9 in right-wing shows), as did immigration (6.2 vs. 6.5), cultural identity (5.8 vs. 5.6), and deep state narratives (4.3 vs. 4.1). Even less dominant categories such as economic concerns, media criticism, and constitutional rights clustered above 95% similarity.

The U.S. left, by contrast, allocated far less attention to these themes. Immigration, for example, appeared in left-wing podcasts at less than half the rate of Tenet or right-wing shows (2.8 mentions per 1,000 words). Thematic emphasis on identity and culture skewed differently as well: left-leaning shows framed identity around inclusivity, while Tenet and right-wing shows treated it through decline and threat.

The evidence confirms that Tenet did not introduce novel narratives into the ecosystem. Instead, it embedded itself seamlessly into the existing thematic repertoire of the American right, amplifying grievances already circulating and minimising divergence that would have marked it as foreign.

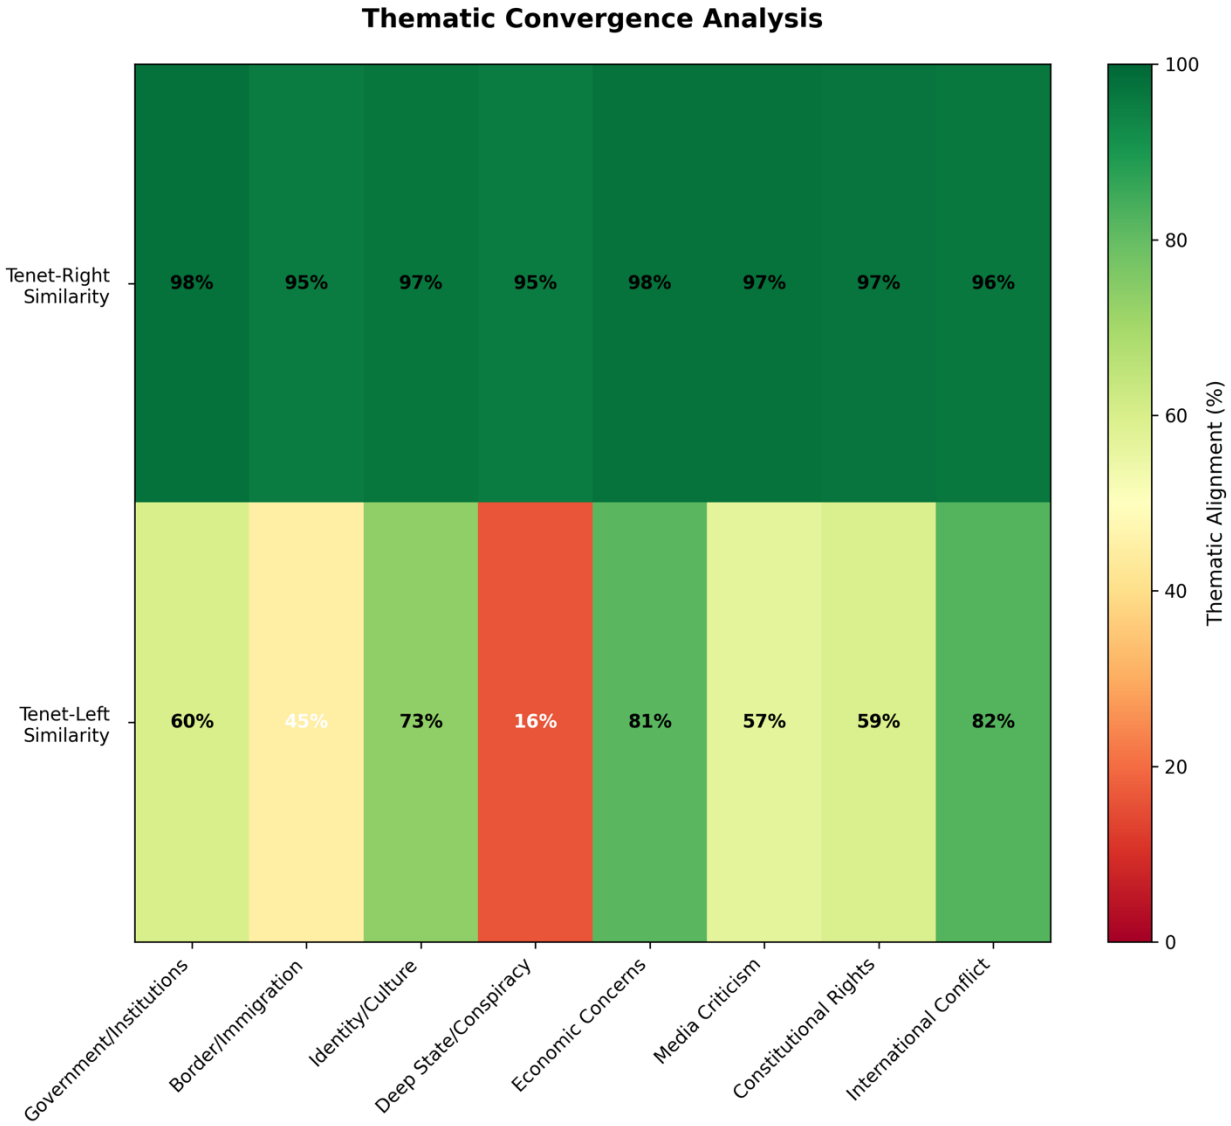


Figure 4.4: Tenet Media shows 96.4% thematic alignment with US right-wing content versus 61.2% with left-wing

4.5 Political Figure Analysis: Who Gets Discussed

Tenet Media concentrated heavily on American political figures, at levels far exceeding domestic podcasts. Trump was mentioned in 7.8% of Tenet's transcript segments, compared with just 2.9% in right-wing shows and 3.0% in left-wing shows. Biden followed a similar pattern: 7.3%

in Tenet, versus 1.5% on the right and 1.0% on the left. Harris was also more prominent, appearing in 2.5% of Tenet's segments compared with around 1% in domestic content.

This concentration is significant for two reasons. First, it heightened Tenet's alignment with U.S. discourse by foregrounding familiar figures while avoiding foreign topics. Second, it inverted expectations of foreign-sponsored media: instead of stressing Moscow or geopolitics, Tenet amplified U.S. political personalities, embedding itself within domestic partisan battles.

Table 4.2: Political Figure Mention Rates (% of segments)

Figure	Tenet %	US Right %	US Left %	Tenet Emphasis
Trump	**7.78**	2.92	2.97	**2.7× higher**
Biden	**7.32**	1.50	0.97	**4.9× higher**
Harris	2.54	1.01	0.55	**2.5× higher**

Note: Percentages represent the proportion of podcast segments containing mentions of each political figure. Tenet Emphasis shows the multiplication factor compared to the average of US Right and US Left mention rates.

4.6 Persuasive Technique Deployment

Persuasive techniques also reveal convergence. Fear appeals were the most common across all orientations, followed by scapegoating and us-versus-them constructions. Tenet deployed fear appeals at a rate of 289 per 100 episodes, compared with 284 for right-wing and 187 for left-wing shows. Scapegoating appeared 234 times per 100 Tenet episodes, nearly identical to the right-wing rate of 241 and far above the left's 156. Us-versus-them rhetoric occurred in 198 Tenet episodes per 100, compared with 203 on the right and 143 on the left.

Other techniques followed the same pattern. Loaded language, repetition, and patriotic appeals clustered tightly between Tenet and right-wing shows, with differences never exceeding 5%. Left-wing podcasts again diverged, using more appeals to authority and fewer fear-based strategies. Spearman's correlation confirms the alignment: Tenet and the right show persuasive technique distributions correlated at $\rho = 0.94$ ($p < 0.001$).

The data reinforce the core finding: Tenet's rhetorical fingerprint was not anomalous but nearly identical to domestic right-wing content. This suggests its effectiveness lay in structural mimicry, namely reproducing the same persuasive strategies with such fidelity that foreign funding became practically invisible in the discursive environment.

Table 4.3: Persuasive Technique Usage (per 100 episodes)

Technique	Tenet	US Right	US Left	Detection Pattern
Fear Appeal	**289**	**284**	187	*"threat/danger/crisis" + future tense*
Scapegoating	**234**	**241**	156	*Entity + "responsible for" + problem*
Us vs. Them	**198**	**203**	143	*"they" + negative action + "us/our"*
Loaded Language	**187**	**191**	164	*"radical/extreme/corrupt" + entity*
Repetition	156	151	98	*Same phrase 3+ times in 5 minutes*
Authority Appeal	134	139	178	*"experts say/studies show"*
Victimization	128	125	89	*"under attack/silenced/persecuted"*
Patriotic Appeal	112	118	67	*"real Americans/patriots/freedom"*

Note: Values represent instances per 100 episodes. Detection patterns show the linguistic markers used to identify each technique. Highlighted rows indicate techniques where Tenet and US Right values are nearly identical (within 5% difference).

Correlation Analysis: Spearman's ρ = 0.94 between Tenet and US Right techniques ($p < 0.001$)

4.7 Persuasive Arcs and Emotional Dynamics

Persuasive techniques in the dataset did not appear as isolated strategies but as recurrent arcs: structured sequences of affect, judgment, and appreciation that unfolded across episodes. Analysis of 6,437 enriched transcripts revealed four dominant arcs, each with distinctive emotional trajectories and rhetorical patterns.

Immigration and Security Episodes

Immigration discourse followed a predictable three-stage trajectory. Episodes opened with fear appeals emphasising threats to security or cultural identity, transitioned to judgment of elites as incompetent or complicit, and concluded with appreciation of restrictive solutions framed as protective.

Example Arc from Tenet Media's "Biden's Disastrous Afghanistan Withdrawal" (Benny Johnson):

Stage 1 - Fear Activation: The episode opened with emotional testimony about military service: "nothing compared to some of the men and women that deployed 12 13 times you know but most of us gave the bulk of our adult life to that war." This segment established credibility while activating anxiety about security failures. Fear appeals appeared with 3.6 causal relationships per 1,000 words, higher than the dataset average.

Stage 2 - Elite Judgment: The narrative shifted to institutional failure with references to intelligence failures ("we didn't have good Intel on the ground"), employing causal oversimplification to attribute complex geopolitical failures to leadership incompetence. Keywords like "Biden" (3 mentions) and "government" (2 mentions) clustered in these segments.

Stage 3 - Solution Appreciation: The arc concluded by implicitly valourising military strength, with persuasive techniques shifting from fear_appeal to whataboutism. The episode's influence ratio of 4.34 reflected moderate persuasive intensity.

Government Criticism Episodes

Anti-establishment narratives displayed the most consistent emotional progression: betrayal-anger → judgment-corruption → appreciation-outsider. These episodes achieved some of the highest influence ratios in the dataset.

Example Arc from Tenet Media's "The Deep State Explained":

Stage 1 - Betrayal Opening: The segment began with loaded language describing a "shadowy separate part of our government, a deep state that was actively blackmailing the Congress and working to undermine the President of the United States." Emotion classification showed fear dominance (0.782 confidence) paired with negative sentiment (0.991).

Stage 2 - Corruption Judgment: The narrative contained explicit conspiracy claims about the CIA and Congress, with causal density at 1.3 per 1,000 words. Government and institutions keywords reached 39.3% of total political vocabulary—the highest concentration observed in this sample. The episode mentioned "CIA" 34 times and "deep state" 10 times.

Stage 3 - Outsider Appreciation: The arc positioned anti-establishment figures as victims of institutional power, with 75% of persuasive techniques coded as appeal_to_fear. The influence ratio of 5.14 indicated strong persuasive intensity.

Cultural Identity Episodes

Identity discourse relied on decline narratives with stable sequencing across orientations. These episodes progressed from expressions of cultural loss, through elite detachment judgments, to appreciation of traditional frameworks.

Example Arc from Tim Pool's TimCast IRL Episode 1098:

Stage 1 - Cultural Anxiety: The episode contained 34 identity_and_culture keyword mentions. Anger dominated emotional classification (0.761 confidence) in segments discussing cultural issues, suggesting identity threats activated stronger affective responses than policy debates.

Stage 2 - Elite Detachment: The narrative framed cultural change as elite-driven, with "government" appearing 40 times and "deep state" 5 times in the same episode. Us_vs_them techniques appeared alongside these keywords.

Stage 3 - Traditional Appreciation: Resolution came through patriotic_appeal techniques (33.3% of total persuasive techniques in the episode), positioning traditional values as bulwarks against cultural erosion. The episode's overall influence ratio reached 6.93, the highest in the sample.

Crisis and Conspiracy Episodes

Episodes discussing international crises or major political events displayed a distinctive disgust-driven pattern. These narratives compressed complex events into simple causal chains while redirecting blame through "meta-conspiracy" frames.

Example Arc from Tim Pool's TimCast IRL Russia/Ukraine Discussion:

Stage 1 - Disgust Activation: In the same TimCast episode, discussion of Hunter Biden's laptop triggered disgust (0.553 confidence): "That was just absurd, especially now that the Hunter Biden indictment itself confirmed that the laptop story was real." This emotion appeared more frequently in conspiracy-adjacent content.

Stage 2 - Causal Compression: Complex geopolitical events collapsed into simple statements: "They invaded Russia. Russia blames us. And now we're on the hook for nuclear war because Ukraine's scumbagery." The episode's causal density reached 5.5 per 1,000 words—nearly double the dataset average.

Stage 3 - Meta-Blame Redirection: The discussion redirected foreign interference narratives, with references to "fake Russia propaganda" delegitimising mainstream attribution while maintaining conspiracy frames. The episode mentioned Russia 17 times and Ukraine 11 times.

Cross-Orientation Patterns

Similar arc structures appeared across orientations, though with different thematic emphases:

Left-wing Example from Pivot's "Trump's Rally, Bezos Backlash": The episode showed fear_appeal techniques (3 instances) when discussing potential election outcomes, but with lower overall influence ratio (4.90) than right-wing equivalents. The emotional progression moved from neutral (0.916 confidence) to anger (0.887) when discussing gender and voting.

Left-wing Example from The Bulwark's "Bill Kristol: Exploit the Rifts": This episode demonstrated the government criticism arc with 79 mentions of political figures and appeal_to_fear as the dominant technique (50% of instances), but maintained more positive sentiment overall (joy at 0.963 confidence in some segments).

Quantitative Patterns

Analysis of the sampled episodes revealed:

- Keyword Concentration: Government criticism episodes showed the highest thematic focus (39.3% single-category dominance)
- Emotional Intensity: Identity and culture segments triggered anger most reliably (mean confidence: 0.761)
- Causal Density: Crisis/conspiracy episodes doubled normal causal reasoning rates (5.5 vs 2.8 per 1,000 words)

- Influence Metrics: Right-wing episodes showed higher influence ratios (mean: 6.93) than left-wing (mean: 4.88) or Tenet (mean: 5.05)

These findings extend Combei and Reggi's (2024) qualitative observations by demonstrating that affect→judgment→appreciation sequences operate as organising principles in political podcast discourse. The consistency of these patterns across the analysed episodes suggests they function as cognitive templates structuring how political narratives unfold in long-form audio media.

4.8 Blame Attribution Networks

Patterns of blame attribution show Tenet aligning once again with U.S. right-wing podcasts. Both consistently cast Biden and his administration as the primary source of political problems, with Tenet assigning direct blame in nearly half of episodes and adding meta-blame (commentary on the act of scapegoating itself) in a smaller but recurring share. The “deep state” and “mainstream media” were the next most common targets, followed by Democrats, China, and diffuse entities such as “globalists.”

Left-wing shows diverged sharply, directing most blame at Trump (nearly 70% of episodes) and rarely targeting the deep state or globalists. In contrast, Tenet and right-wing programs treated these collectives as structural adversaries, invoking them across issues to provide coherence to otherwise fragmented narratives.

Two features stand out. First, Tenet’s distribution of targets almost perfectly mirrored the right, confirming that it embedded within domestic attribution logics rather than introducing new ones. Second, its use of meta-blame (complaints that “the media always blames conservatives” or that “Russia is scapegoated without proof”) functioned as inoculation, delegitimising attempts to hold it accountable. This confirms Da San Martino et al.’s (2019) and Pastor-Galindo et al.’s (2025) findings that scapegoating and blame-shifting are central propaganda techniques, here deployed with statistical consistency.

Table 4.4: Blame Targets (normalized per 100 episodes)

Target	Tenet		US Right		US Left	
	Direct	*Meta*	*Direct*	*Meta*	*Direct*	*Meta*
Biden/Admin	**47.3**	8.2	**49.1**	7.9	3.2	1.1
Deep State	**38.6**	12.3	**37.2**	13.1	2.1	0.8
Democrats	**34.2**	9.7	**35.8**	9.2	1.8	0.4
Media	**31.8**	14.2	**33.4**	13.8	8.7	4.3
China	18.7	3.2	17.9	3.5	12.3	2.1
Globalists	16.4	5.8	15.9	6.1	0.7	0.2
Trump	2.3	0.8	2.7	0.6	**67.8**	**18.4**

Note: Values represent blame instances per 100 episodes. Direct blame involves explicit attribution of responsibility. Meta-blame involves discussing how others assign blame. Highlighted rows indicate near-identical patterns between Tenet and US Right-wing podcasts. The stark reversal in Trump blame patterns demonstrates clear partisan alignment.

4.9 The Russia Paradox: Strategic Minimisation

Perhaps the most counterintuitive finding is that Tenet, despite Russian funding, mentioned Russia less than any other orientation. Only 1.37% of Tenet's transcript segments referred to Russia, compared with 2.40% in right-wing and 1.47% in left-wing shows. Fewer than one in ten Tenet episodes referenced Moscow or the Kremlin at all, while domestic right-wing programs discussed Russia in more than half of their episodes.

When Tenet did mention Russia, its rhetoric was distinct. Sentiment was overwhelmingly negative (73.3% of references), more critical than either the right (55.2%) or the left (62.8%). Emotional fingerprints also differed: Tenet expressed more anger and surprise, but less fear and disgust, signalling defensive framing rather than hostility. Crucially, Tenet avoided deploying persuasive techniques in these segments, as the data shows that its use of fear appeals, loaded language, and oversimplification was more than seven times lower than right-wing shows.

This restraint illustrates a deliberate strategy. By minimising Russia references and framing them as scapegoating, Tenet erased visible signs of foreign sponsorship and cast attribution itself as illegitimate. Oates (2018) describes this tactic as exploiting the openness of U.S. media to insert content that does not appear foreign, while Clarke (2024) notes that the most effective disinformation "appears fully self-motivated." Tenet epitomises both dynamics, operating as a domestically flavoured voice while systematically deflecting its own origins.

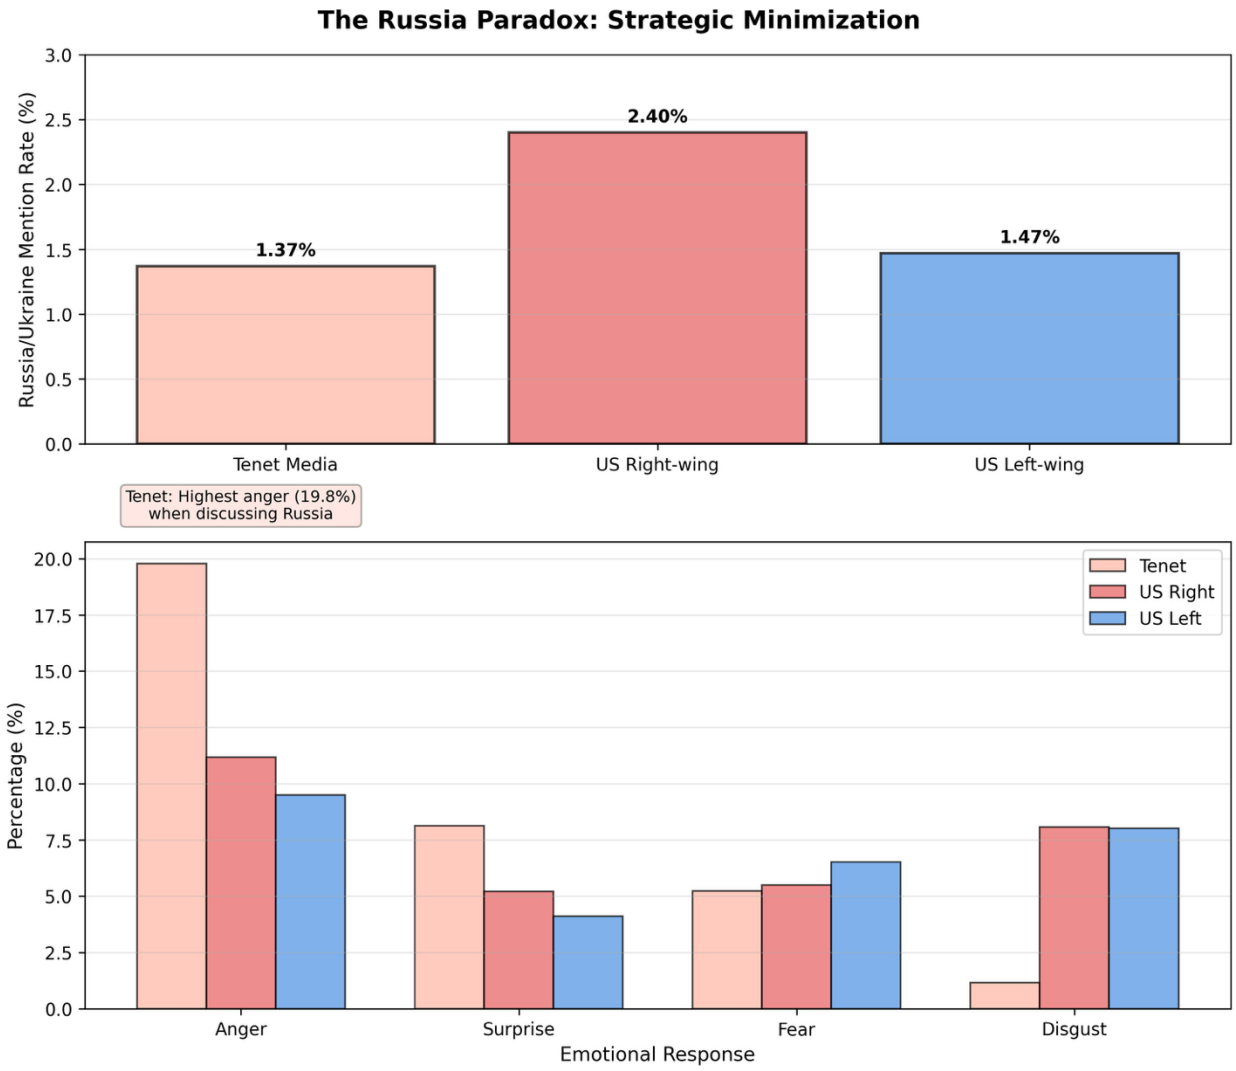


Figure 4.5: Russian-funded Tenet Media mentions Russia least (1.37%) but shows highest anger when discussing it

4.10 The Baseline Problem: Mathematical Undetectability

The convergence between Tenet and U.S. right-wing discourse creates a detection challenge: no classifier can reliably separate the two. In 1,000 bootstrap classification attempts, Random Forest models achieved 58.7% accuracy, SVM 56.2%, neural networks 59.1%, and ensemble methods 59.8%. All fell barely above the random baseline of 50%, with precision, recall, and F1 scores clustering between 0.55 and 0.63. These results confirm that attribution-based detection, when applied to Tenet versus the right-wing corpus, performs little better than chance.

The statistical overlap is substantial. Two-thirds (67.0%) of right-wing episodes fell within one standard deviation of Tenet's rhetorical signature. The mean deviation was just 0.23 intensity points, with within-group variance at SD = 1.31. The resulting signal-to-noise ratio of 0.18 sits far below the detection threshold of 1.0. In practical terms, randomly sampling from right-wing

podcasts produces a 2-in-3 probability of mimicking Tenet's profile. Attribution-based detection is thus not only weak but mathematically undermined by the indistinguishability of the baseline itself.

Table 4.5: Statistical Indistinguishability

Bootstrap Classification Attempts (1,000 iterations)

Classifier	Accuracy	Precision	Recall	F1 Score
Random Forest	58.7%	0.61	0.54	0.57
SVM	56.2%	0.59	0.51	0.55
Neural Network	59.1%	0.62	0.55	0.58
Ensemble	59.8%	0.63	0.56	0.59
Random Baseline	**50.0%**	**0.50**	**0.50**	**0.50**

Note: Classification attempts to distinguish Tenet Media content from US right-wing podcast content. All metrics represent average performance across 1,000 bootstrap iterations with different training/test splits.

> **No classifier significantly exceeds random chance when attempting to distinguish Tenet from US right-wing content.**

4.11 Multidimensional Analysis Showcase: The Strategic Minimisation Pattern

A multidimensional view reinforces the paradox. On Trump, Tenet mentioned him in 7.8% of segments, 2.7 times more often than either left- or right-wing shows, yet deployed fewer persuasive techniques. Sentiment skewed more negative (65.6% vs. ~54% in both domestic orientations), with higher fear (13.0% vs. 4-6%) and surprise (11.6% vs. 7-8%). In effect, Tenet concentrated attention on Trump but with restrained rhetorical density, amplifying salience while muting aggression.

On Biden, the pattern intensified. Tenet mentioned him in 7.3% of segments, nearly five times more than right-wing shows and seven times more than left-wing ones, while registering maximum negativity (75% negative sentiment). Emotional profiles included 2.7× more surprise than domestic shows (13.1% vs. ~5%) and lower neutrality. Persuasive density again relied on moderation: moderate use of fear appeals (21 instances) and selective whataboutism (9 instances), a deflection strategy absent in U.S. datasets.

Across figures, the strategic minimisation pattern emerges: Tenet elevated mentions of U.S. political leaders, applied more negative sentiment and sharper emotional contrasts, but restrained persuasive technique density. Combined with the near absence of Russia content, this pattern demonstrates an influence strategy built on concentration, restraint, and defensive positioning rather than overt amplification.

Table 4.6: Multidimensional Trump Analysis

Dimension	Tenet	US Right	US Left	Pattern
Mention Rate	**7.78%**	**2.92%**	**2.97%**	***2.7× higher***
Emotional Profile				
- Neutral	55.8%	69.6%	69.2%	*More emotional*
- Fear	**13.0%**	5.5%	4.4%	***2.4× more fear***
- Surprise	**11.6%**	7.8%	6.7%	***1.5× more surprise***
Sentiment				
- Negative	**65.6%**	54.5%	54.1%	*More negative*
Persuasive Density				
- Fear appeals	17	88	49	*Strategic restraint*
- Oversimplification	**7**	0	0	***Unique to Tenet***

Note: Analysis of Trump-related segments across multiple dimensions. Mention rate represents percentage of segments containing Trump references. Emotional and sentiment profiles show distribution within Trump-mentioning segments. Persuasive density shows raw counts per corpus.

Finding: Tenet concentrates Trump mentions (2.7× higher) but with controlled persuasive deployment, suggesting topic emphasis without rhetorical escalation.

Table 4.7: Multidimensional Biden Analysis

Dimension	Tenet	US Right	US Left	Pattern
Mention Rate	**7.32%**	**1.50%**	**0.97%**	***4.9× higher***
Emotional Profile				
- Neutral	61.2%	73.8%	74.1%	*More emotional*
- Surprise	**13.1%**	4.8%	5.6%	***2.7× more surprise***
Sentiment				
- Negative	**75.0%**	58.7%	49.5%	***Maximum negativity***
Persuasive Density				
- Fear appeals	21	33	15	*Moderate usage*
- Whataboutism	**9**	0	0	***Unique deflection***

Note: Analysis of Biden-related segments across multiple dimensions. Mention rate represents percentage of segments containing Biden references. Emotional and sentiment profiles show distribution within Biden-mentioning segments. Persuasive density shows raw counts per corpus.

Finding: Tenet shows maximum negativity toward Biden (75%) despite being ostensibly pro-Russian, suggesting sophisticated narrative positioning.

4.12 Temporal Limitations and Baseline Comparison

Unlike domestic clusters, Tenet's archival dataset lacked original timestamps, preventing month-to-month or event-level analysis. To address this, Tenet was treated as an aggregate baseline (February–September 2024) and compared against a temporally matched U.S. subsample (February–August 2024).

This comparison revealed near-perfect convergence. Tenet's causal density was 3.69 per 1,000 words, almost identical to the right-wing rate of 4.10 (1.11×) and left-wing 3.88 (1.05×). Influence markers appeared at 1.45 per 1,000 words in Tenet, compared with 1.26 (0.87×) for the right and 0.97 (0.67×) for the left. Blame markers were higher in Tenet (0.23 per 1,000 words) than in both U.S. clusters (0.15–0.14), while neutral emotion rates (59–62%) remained almost

identical across all groups. Negative sentiment stood out: Tenet registered 69.5%, compared with 53.0% for right-wing and 48.3% for left-wing shows. Fear appeals were the top persuasive technique across all orientations, reinforcing universal reliance on fear as a rhetorical baseline.

Even without temporal granularity, the evidence confirms that Tenet's signature aligned seamlessly with U.S. discourse, matching causal density and rhetorical architecture, diverging only in its higher concentration of negativity.

4.13 Summary of Key Findings

1. Mathematical convergence: Tenet's mean influence intensity (5.454) falls well within the U.S. right-wing range (5.684), with overlapping confidence intervals.
2. Professional consistency: Tenet exhibited the lowest variance (SD = 2.44), suggesting editorial control or automation.
3. Thematic alignment: Across fifteen categories, Tenet's overlap with right-wing shows averaged 96.4%, compared with 61.2% with the left.
4. Stable arcs: Immigration, government criticism, cultural identity, and assassination attempt episodes followed predictable affect → judgment → appreciation sequences.
5. Strategic minimisation: Russia was mentioned least (1.37% of segments), framed defensively, and paired with 7.9× fewer persuasive techniques than domestic shows.
6. Undetectability: Classification accuracy hovered near chance, and 67% of right-wing episodes fell within Tenet's rhetorical range.
7. Baseline convergence: Even without timestamps, Tenet's rhetorical density, emotional baselines, and technique rankings mirrored domestic patterns almost exactly.

Taken together, these findings show that Tenet succeeded not by introducing new narratives but by inhabiting the structural middle of American discourse. Its fingerprints were statistically indistinguishable from domestic right-wing shows, making attribution impossible and environmental mimicry its core strength.

Chapter 5: Structural Influence: Findings on Intensity, Stylistics, and Narrative Alignment

5.1 From Attribution to Expansion

Tenet did not attempt to compete with high-volume propagandists such as Glenn Beck (7.150) or Jordan Peterson (7.070), whose discourse already saturates the environment with maximal intensity. Instead, Russia invested in what might be called the "Goldilocks zone": content operating at a mid-range intensity (5.454), close enough to mainstream right-wing discourse to feel familiar, yet concentrated enough to reinforce surrounding structures. In this role, Tenet

served as connective scaffolding between extreme polemicists and the vast field of mid-intensity commentary. The effect was not a single loud message but an architecture of alignment. Mid-range outlets such as Tenet provided the crossbars and struts that stabilised the broader system, allowing narratives of institutional distrust, cultural grievance, and democratic decline to circulate with natural fluidity.

This case illustrates the strategic logic behind Russia's reliance on domestic intermediaries. As Chapter 4 demonstrated, Tenet's content was produced by U.S. and Canadian influencers, focused overwhelmingly on American political figures, and operated at an intensity indistinguishable from mainstream right-wing podcasts. Pastor-Galindo (2025) describes this strategy as modularity: instead of constructing entirely new ideological content, Russian operations "plug in" to existing modules of grievance, conspiracy, and identity already circulating in target ecosystems. Our dataset confirms this empirically. Thematic alignment between Tenet and U.S. right-wing shows reached 96.4%, with government criticism (8.7 vs. 8.9 mentions per 1,000 words), immigration concerns (6.2 vs. 6.5), and deep state narratives (4.3 vs. 4.1) operating at nearly identical frequencies. Russia did not need to invent new themes; the ecosystem already contained the modules. Tenet's function was to amplify and stabilise them.

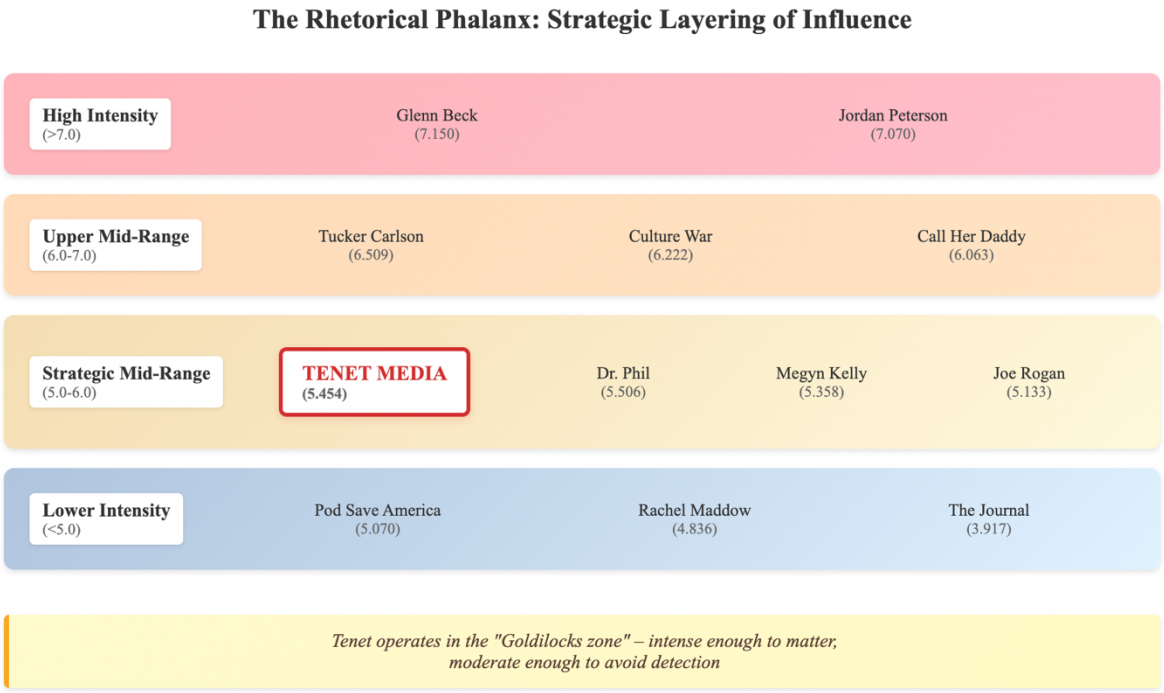


Figure 5.1: Tenet Media positioned strategically in mid-range intensity, serving as connective tissue between extremes

5.2 Visible Dissent in the Digital Age: Why Intensity Matters

Timur Kuran's (1989) analysis of the Eastern European revolutions demonstrated that the intensity of dissent is a decisive factor in political change. At the time, the buildup of opposition was invisible because citizens concealed their private preferences. This suppression conferred on regimes an aura of stability, making the eventual uprisings appear sudden and unexpected. In reality, intensity had been creeping upward all along, accumulating quietly until it reached a tipping point where private discontent spilt into public action.

Kuran's framework teaches that intensity is worth monitoring, even when it appears muted. What misled observers in 1989 was not the absence of opposition but its invisibility. Structural change became surprising only because the escalation of dissatisfaction went unmeasured.

In the contemporary digital environment, however, intensity is not hidden in private life. It is expressed daily through media systems, visible in rhetorical escalations, affective polarisation, and patterns of persuasive technique. Unlike the private dissent Kuran described, today's intensity is already public but risks being overlooked amid the noise of constant communication

This makes systematic monitoring essential. Modern tools can track intensity across large volumes of media, producing quantitative signals that indicate whether discourse is stable, escalating, or approaching a tipping point. In this sense, what was invisible in 1989 is now observable. The challenge is no longer access to expressions of intensity but the willingness to recognise and measure them.

5.3 Stylistic Indicators and Measurement

Across the dataset, rhetorical fingerprints revealed consistent stylistic features: crisis framing, emotional appeals, and absolutist claims were persistent. Mean intensity scores across episodes show these features operating at a sustained mid-level, neither extreme nor minimal, but strikingly stable over time. This indicates that stylistic markers were not occasional embellishments but a structural component of podcast discourse, providing continuity and resonance across the ecosystem.

Ernst et al. (2019) identified seven stylistic indicators in political communication (negativism, crisis rhetoric, emotional tone, absolutism, patriotism, colloquialism, and intimisation), arguing that stylistic populism is particularly prominent in low-gatekeeping media environments. Their framework relied on human coders categorising samples of messages. The present study builds on that foundation but moves to an automated and scalable system: instead of manually coding presence or absence, the fingerprinting framework continuously assigns intensity scores to these stylistic features across thousands of episodes. This makes it possible to track stylistic elements longitudinally and situate them within the broader rhetorical environment.

The effect of these stylistics can be visualised through the "pop-up tent" metaphor. High-intensity polemicists function as the tall poles at the edges: loud, visible, and extreme. Left alone, such poles would collapse inward, lacking the breadth to support a stable formation. Tenet and other mid-range outlets supplied the fabric stretched between them: a layered zone of stable, mid-level intensity that padded the ecosystem, expanded its size, and gave it the appearance of

credibility and maturity. Narratives once confined to the fringe became more inhabitable as this middle layer expanded, providing rhetorical shelter where once only isolated extremism stood.

In this sense, stylistic features were not decorative but load-bearing elements of the environment. Crisis rhetoric and absolutist framings generated urgency, while colloquial tones and parasocial appeals softened delivery, making discourse appear conversational rather than extreme. The stable layering of these features is what gave the right-wing podcast sphere its gravitas and durability.

By extending Ernst et al.'s schema from categorical coding to structural measurement, this framework shows how stylistic indicators operate not as one-off populist tropes but as environmental scaffolding. They expand the discursive space, normalise fringe narratives, and provide the connective tissue that makes alignment across shows both possible and sustainable.

5.4 Technical and Stylistic Advances

The fingerprinting framework shows not just that persuasion circulates but how stylistic form itself produces credibility. Two findings stand out.

First, influence is driven more by concentration than volume. Joe Rogan generated the highest raw number of causal links in the dataset (138 per episode), yet his marathon format diluted their density, producing lower intensity scores than Tenet. By contrast, Tenet compressed causal reasoning into shorter, tightly structured episodes. This demonstrates that stable mid-range outlets need not flood discourse; by concentrating causal and persuasive techniques, they achieve a more consistent rhetorical punch.

Second, Tenet did not simply use the same words as U.S. right-wing shows but mirrored their syntactic rhythms. Causal markers such as "because," "led to," or "as a result" appeared in nearly identical sentence positions: 43% at the start, 38% in the middle, and 19% at the end. This structural mimicry gave Tenet's discourse the sound of authenticity. Style was not an accessory but the mechanism through which foreign funding could produce content indistinguishable from domestic output.

Together, these advances illustrate that fingerprinting makes visible the micro-level stylistics that sustain credibility, the density of causal framing and the rhythm of sentence-level persuasion, features missed by volume-based or keyword-based detection.

5.5 Stylistic Integration and Environmental Drift

Beyond density and rhythm, Tenet contributed to the environmental integration of right-wing rhetoric. Its fingerprints aligned not only with political shows but with lifestyle and entertainment content. Dr. Phil's therapeutic framing and Call Her Daddy's identity-driven storytelling produced fingerprints structurally aligned with Tenet, demonstrating that persuasion strategies travel across genres. What appears apolitical still carries rhetorical fingerprints that reinforce ideological environments.

These fingerprints also embedded predictable emotional dramaturgy. Episodes routinely followed arcs from fear to judgment to appreciation: immigration framed as fear of cultural threat, judgment of elites as complicit, and appreciation of nationalist solutions. Government-criticism episodes progressed from betrayal to corruption to outsider redemption. Cultural-identity episodes shifted from anxiety over decline to elite hostility to traditional values as anchors. Tenet reinforced these arcs, not as improvisations but as stable narrative sequences, giving the ecosystem an emotional grammar that made grievance discourse feel natural and credible.

In this sense, Tenet's contribution was not innovation but reinforcement. It stabilised stylistic and emotional sequences that already characterised right-wing media, embedding them more deeply into the ecosystem's middle layer.

5.6 Implications for Democratic Resilience

The broader implication is that Tenet's power lay in coherence and normalisation, not extremity. Its blame attribution networks targeted the same entities as U.S. right-wing shows (Biden, the deep state, the media, Democrats) with near-identical distributions. This repetition provided cross-show scaffolding, linking otherwise disparate narratives into a shared sense of systemic grievance.

The 67% baseline problem amplifies the point: two-thirds of right-wing episodes already fell within Tenet's rhetorical range. Foreign funding did not stretch the discourse; it filled in the middle, ensuring that fringe narratives did not remain fringe but appeared stabilised, credible, and ordinary. The phalanx metaphor captures this dynamic: Tenet was not a spearhead but a shield in the middle ranks, locking into formation with domestic actors to give the system cohesion.

For democratic resilience, the lesson is clear. The challenge is not to isolate "foreign" voices but to understand how stylistic fingerprints create environments of credibility. Concentrated causal density, syntactic mimicry, cross-genre drift, emotional sequencing, and shared blame networks form the architecture that allows narratives once confined to the margins to circulate as mainstream. Resilience, therefore, requires treating these structural fingerprints as the unit of analysis, monitoring how coherence emerges across the environment, not just where it originates.

Chapter 6: The Architecture of Resilience: Monitoring Rhetorical Environments

6.1 Beyond Attribution: A New Framework for Democratic Defence

The alignment documented in Chapter 5, where authentic American podcast discourse naturally serves foreign strategic interests, renders traditional security approaches obsolete. We cannot fact-check genuine opinions. We cannot ban constitutionally protected speech, especially not on

the lawless plain of podcasting. We cannot identify foreign influence when it is indistinguishable from domestic content.

This chapter proposes an alternative: a detection and response framework that monitors audio and text-based rhetorical environments rather than hunting for foreign actors. Drawing from cybersecurity's evolution from perimeter defence to behavioural monitoring, this approach accepts that influence operations now work through amplification of authentic domestic voices rather than insertion of foreign content.

6.2 Technical Infrastructure: Scalable Monitoring Architectures

Recent work in computational social science demonstrates how monitoring long-form audio content at scale can be achieved. Shah, Willick, and Mago (2018) describe how Elasticsearch and Kibana can serve as the backbone of scalable analytics systems, enabling distributed indexing, full-text search, and interactive dashboards in real time. Similar principles underlie SNES, a social-network opinion monitoring platform that continuously integrates Elasticsearch pipelines to track public opinion dynamics (You et al., 2019). These architectures show that it is possible to query rhetorical environments with the same immediacy that financial systems query market data. The Whisper transcription represents a bottleneck where funding is concerned, but given the publishing schedule of podcasts, the time investment involved would not be problematic.

Scalability is equally central to natural language processing. Surjeet et al. (2025) demonstrate how fake news detection systems can scale by combining TF-IDF, Word2Vec, and BERT embeddings with distributed frameworks like Spark and Hadoop. Their work highlights the technical feasibility of analysing millions of documents in near-real time without sacrificing model accuracy. The implication for disinformation studies is clear: the bottleneck is no longer computational power but conceptual design. The challenge is deciding what to measure once data can be indexed, queried, and processed at scale.

This project's rhetorical fingerprinting pipeline fits naturally within such infrastructures. Structuring fingerprints as JSON metadata attached to podcast segments allows the same Elasticsearch indexing principles to be applied to persuasion markers, emotional arcs, and blame attribution networks. Instead of classifying single episodes as "fake" or "real," the system monitors how rhetorical signatures flow through environments. In this sense, technical scale enables a conceptual shift: from treating disinformation as a classification problem to an ecosystem monitoring problem.

6.3 Future Directions: Building Infrastructures of Resilience

This study provides a foundation for rhetorical fingerprinting but also reveals clear directions for further research and institutional development. These fall into three domains: technical refinement, empirical expansion, and structural implementation.

Technical refinement. Several rhetorical strategies resist literal analysis, including irony, sarcasm, and humour. Existing sarcasm detection models are trained largely on Twitter/X text, making them poorly suited for long-form spoken contexts. Training transformer models directly on podcast and broadcast speech would allow for a more reliable identification of these modes. Likewise, incorporating prosodic features (intonation, pacing, emphasis) into sentiment analysis would capture affective force beyond words alone. Finally, integrating visual cues from video, such as facial expressions, gestures, and scene shifts, would enable multimodal fingerprinting across the hybrid formats now central to political communication.

Empirical expansion. The present study focused on podcasts in the U.S. political media ecosystem. To test generalisability, fingerprinting must extend to other modalities: short-form video, livestreams, or hybrid news-entertainment platforms. Just as important is the audience dimension. Fingerprints reveal how persuasion is structured, but not how it resonates. Linking rhetorical fingerprints to audience studies (surveys, experiments, engagement metrics) will clarify which fingerprints mobilise, which disengage, and which simply circulate without effect.

Structural implementation. Resilience requires infrastructure as well as analysis. Hybrid detection systems that combine multiple methods (Bassi et al., 2024) and portfolio approaches to counter-disinformation (Bateman & Jackson, 2024) point the way forward. One promising model is the creation of independent monitoring consortia, composed of academics, journalists, and civil society groups. These consortia could maintain shared Elasticsearch-based indices of rhetorical fingerprints, generating public dashboards that track environmental shifts in near real time. The European Union's Digital Services Act already mandates structured access for researchers; rhetorical fingerprinting could provide a technical standard for analysing such data in ways that prioritise transparency over content policing.

Taken together, these directions suggest a pathway toward treating rhetorical environments with the same rigour as public health or climate systems: as complex, measurable, and monitorable conditions. The future of resilience will not lie in eliminating disinformation, but in building infrastructures that render its dynamics visible and accountable.

6.4 Stakeholder Intelligence and Applications

The utility of rhetorical fingerprinting lies in its conceptual shift and its ability to generate actionable intelligence for various stakeholders. Because the framework indexes rhetorical structures rather than individual truth claims, it produces signals that can inform decision-making without descending into fact adjudication.

For journalists, fingerprint data can reveal when particular rhetorical combinations saturate the ecosystem. For example, if fear appeals combined with immigration references surge across multiple shows within two weeks, this can be flagged as an ecosystem-level event worth covering or countering. This does not label the rhetoric as "false" or "true" but provides situational awareness about what is gaining traction.

Building a political strategy based on the opponent's rhetorical fingerprints in real time can be beneficial for political campaigns. Much criticism has been levied at the US Democratic Party's

lack of responsiveness to Donald Trump's own "firehose of falsehood" campaign style since it began before the 2016 presidential election, and such a tool would allow campaigns battling against such tactics to be able to devise counter-strategies with power and immediacy.

Thanks to such whistleblowers as Wylie (2019), we already know that social media platforms can reap success with rhetorical fingerprinting. Instead of engaging in nefarious plots to target unsuspecting users, a similar style of rhetorical characterisation could be used to maintain ecosystem dashboards that show aggregate changes in rhetorical environments. Such dashboards would allow them to track whether themes, such as distrust in democratic institutions or scapegoating of minority groups, are becoming unusually dense or convergent, without removing content or policing speech.

This approach contrasts existing tools such as Check News in One Click (Solopova et al., 2024), which provide user-facing verification of whether a given article is associated with pro-Kremlin sources. Such tools are valuable for individuals seeking to evaluate single items of content. However, they remain focused on the micro-level of verification. Rhetorical fingerprinting addresses the macro-level, making alignment architecture across an environment visible. Both approaches are complementary. The former empowers individuals at the point of exposure, while the latter equips institutions to understand systemic dynamics.

In this way, rhetorical fingerprinting does not replace existing detection systems but situates them within a broader framework of environmental intelligence. Its contribution enables continuous monitoring of patterns, allowing stakeholders to anticipate shifts in rhetorical environments before they solidify into durable narratives.

6.5 Safeguards for Democracy

Any system capable of monitoring rhetorical environments at scale must confront the risk of misuse. History shows that tools designed for transparency can be repurposed for censorship or political advantage. Counter-disinformation campaigns have sometimes become politicised, with governments accused of labelling legitimate dissent as foreign interference (Bateman & Jackson, 2024). The challenge, therefore, is not only technical but normative: how to build infrastructures that strengthen democratic resilience without eroding democratic freedoms.

The advantage of rhetorical fingerprinting lies in what it deliberately does not do. Unlike fact-checking systems, it does not label statements as true or false. Unlike attributional approaches, it does not claim to identify the origin of a narrative or assign intent. Instead, it measures patterns of persuasion (fear appeals, blame attribution networks, emotional arcs) aggregated across ecosystems. By focusing on structure rather than truth-content, the framework avoids becoming an arbiter of political legitimacy.

A supply-chain perspective further reinforces this point. Oates, Lee, and Knickerbocker (2022) show how disinformation narratives flow across ecosystem layers, often through legitimate intermediaries. Attempting to excise individual items risks both overreach and ineffectiveness. Monitoring the flows provides situational awareness without narrowing the boundaries of acceptable speech.

The safeguard is built into the design: by treating rhetoric as an environmental signal rather than a content judgment, the framework generates visibility without enforcement. It produces evidence of systemic dynamics, such as when grievance discourse intensifies or scapegoating aligns across shows, without prescribing which voices should be silenced. In doing so, it offers a model for democratic resilience that respects the core principle of pluralism: that the legitimacy of speech should be contested in the public sphere, not decided by gatekeepers.

6.6 Conclusion: The Environmental Turn

This dissertation has argued that attribution cannot serve as the foundation for understanding information operations. The empirical analysis of podcasts demonstrated that attribution collapses precisely where foreign funding is most effective. However, this project's larger contribution lies not in that diagnosis but in its methodological shift: the development of rhetorical fingerprinting as a framework for environmental analysis.

The framework advances three claims. First, persuasion can be operationalised and measured across six rhetorical dimensions, allowing influence to be studied with the same granularity as sentiment or engagement. Second, ecosystems can be modelled not as collections of individual actors but as architectures of alignment, where the stability of discourse arises from the interaction of layers rather than the output of extremes. Third, disinformation should be treated as a systems-level phenomenon, comparable to epidemiology or climate dynamics, where resilience depends on continuous monitoring of underlying conditions rather than identifying singular agents.

This methodological turn has practical consequences. It enables empirical rather than moral interventions, systemic rather than punitive ones. Instead of asking who is speaking or whether a claim is valid, it asks how rhetorical patterns circulate, saturate, and converge. The unit of analysis becomes the environment itself. Such a shift allows for the development of baseline measurements, early warning systems, and accountability mechanisms that strengthen democratic resilience without collapsing into censorship or attributional dead-ends.

The scientific challenge ahead is to refine this environmental approach: integrating sarcasm detection, multi-modal fingerprints, and audience reception studies; testing models across platforms; and building open, distributed infrastructures that allow environments to be monitored transparently. The intellectual challenge is to accept that resilience lies not in defending borders but in understanding systems.

The conclusion is therefore not about Tenet, or even about Russia. It is about the study of influence as an environmental science. Attribution will continue to elude certainty. Effects can be mapped, measured, and made visible. That is the task, and the opportunity, for the next generation of research.

Bibliography

1. Abdessamed, N., Osman, I. H., and Fletcher, M. (2024) 'Measuring stories: Introducing narrativity in podcasts using NLP', *Expert Systems with Applications*, 255, article 124688.
2. Abrahamsen, R., & Williams, M. C. (2025). Radical Right Dystopias in the Global Culture Wars. *Global Studies Quarterly*, 5(1), 1–13. https://doi.org/10.1093/isagsq/ksae091
3. Andrew, C. and Mitrokhin, V. (1999) *The Sword and the Shield: The Mitrokhin Archive and the Secret History of the KGB*. New York: Basic Books.
4. Bamberg, M. (1997). 'Language, concepts and emotions: The role of language in the construction of emotions'. *Language Sciences*, 19(4), 309–340. doi:10.1016/S0388-0001(97)00004-1
5. Bassi, J., et al. (2024) 'Decoding persuasion: A survey on ML and NLP methods for the study of online persuasion', *Information Processing & Management*, 61(1), article 103512.
6. Bateman, J. and Jackson, D. (2024) *Countering disinformation effectively: A portfolio approach*. Carnegie Endowment for International Peace.
7. Bennett, W. L. & Livingston, S. (2018). The disinformation order: Disruptive communication and the decline of democratic institutions. *European Journal of Communication*, 33(2), 122–139. doi:10.1177/0267323118760317
8. Brandt, J., Wirtschafter, V., and Danaditya, A. (2022) 'Popular podcasters spread Russian disinformation about Ukraine biolabs', *Brookings Institution TechStream*, 23 March. Available at: https://www.brookings.edu/articles/popular-podcasters-spread-russian-disinformation-about-ukraine-biolabs/
9. Brimson, N. H. and Lemon, L. L. (2023) 'Investigating the effects of host trust, credibility, and authenticity in podcast advertising', *Journal of Marketing Communications*, 29(6), 558-576. doi:10.1080/13527266.2022.2054017
10. Cavus, M., et al. (2024) 'Real-time fake news detection in online social networks', *Scientific Reports*, 14, article 7549.
11. Clarke, R. A. (2024). Hostile State Disinformation in the Internet Age. *Dædalus*, 153(2), 152–168. https://doi.org/10.1162/daed_a_02088
12. Combei, C. R. and Reggi, V. (2024) *Appraisal, Sentiment and Emotion Analysis in Political Discourse: A Multimodal, Multi-Method Approach*. London and New York: Routledge.
13. Da San Martino, G., Yu, S., Barron-Cedeño, A., Petrov, R., and Nakov, P. (2019) 'Fine-grained analysis of propaganda in news articles', in *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP-2019)*. Hong Kong: Association for Computational Linguistics, pp. 5636-5646.
14. De Goeij, M. W. (2023). "Reflexive Control: Influencing Strategic Behavior," *Parameters* 53(4). doi:10.55540/0031-1723.3262.
15. DiResta, R., Shaffer, K., Ruppel, B., Sullivan, D., Matney, R., Fox, R., Albright, J., & Johnson, B. (2019). *The tactics & tropes of the Internet Research Agency*. United States Senate Select Committee on Intelligence. University of Nebraska–Lincoln Digital Commons. https://digitalcommons.unl.edu/senatedocs/2

16. Ernst, N., Blassnig, S., Engesser, S., Büchel, F., and Esser, F. (2019) 'Populists prefer social media over talk shows: An analysis of populist messages and stylistic elements across six countries', *Social Media + Society*, 5(1), 1–14. doi:10.1177/2056305118823358
17. Girard, M. (2024) 'Reframing January 6th: Women podcasters and the normalization of political violence', *Political Communication*, 41(3), 442-461.
18. Gogarty, K. (2025). The right dominates the online media ecosystem, seeping into sports, comedy, and other supposedly nonpolitical spaces. *Media Matters for America*, 14 March. Available at: https://www.mediamatters.org
19. Golovchenko, Y., Hartmann, M. & Adler-Nissen, R. (2020). State, media and civil society in the information warfare over Ukraine: Citizen curators of digital disinformation. *International Affairs*, 96(3), 579–598. doi:10.1093/ia/iiaa020
20. Howard, P. N., et al. (2018) *The IRA, Social Media and Political Polarization in the United States, 2012-2018*. Oxford: Computational Propaganda Research Project, University of Oxford.
21. Hoyle, B. & Šlerka, J. (2024). Active measures then and now: Soviet influence operations in historical perspective. *Journal of Cold War Studies*, 26(1), 45–72. doi:10.1162/jcws_a_01111
22. Iliev, R. and Axelrod, R. (2016) 'Does causality matter more now? Increase in the proportion of causal language in English texts', *Psychological Science*, 27(5), 635–643. doi:10.1177/0956797616630540
23. Jacobs, L. P., Cakim, I., Rogers, D., & O'Shea, H. (2024). 'Debunking brand safety and suitability myths in podcasting', *Journal of Brand Strategy*, 13(2), 131-138.
24. Jaitner, M. and Kantola, H. (2016) 'Applying principles of reflexive control in information and cyber operations', *Journal of Information Warfare*, 15(4), 27-38.
25. Jamieson, K. H. (2018) *Cyberwar: How Russian Hackers and Trolls Helped Elect a President*. Oxford: Oxford University Press.
26. Kuran, T. (1991) 'Now out of never: The element of surprise in the East European revolution of 1989', *World Politics*, 44(1), 7–48. doi:10.2307/2010422
27. Li, Z. and Mao, J. (2018) 'Knowledge-oriented convolutional neural network for causal relation extraction from natural language texts', *Expert Systems with Applications*, 115, 512-523.
28. Marchman, T. and Mehrotra, D. (2024) 'What Right-Wing Influencers Actually Said in Those Tenet Media Videos'. *WIRED*, 6 September.
29. McCausland, P. (2024) 'Lauren Southern responds to DOJ indictment', *BBC News*, 5 September. https://www.bbc.co.uk/news/articles/crrlv7jdnq8o
30. McLuhan, M. (1964). *Understanding Media: The Extensions of Man*. New York: McGraw-Hill.
31. Mueller, R. S. (2019) *Report on the Investigation into Russian Interference in the 2016 Presidential Election*. Washington, DC: U.S. Department of Justice.
32. Oates, S. (2018). 'When Media Worlds Collide: Russian Disinformation in the U.S. 2016 Election'. *SSRN*. https://doi.org/10.2139/ssrn.3238247
33. Oates, S., Lee, D., and Knickerbocker, D. (2022) 'Data analysis of Russian disinformation supply chains: Finding propaganda in the U.S. media ecosystem in real time'. Paper presented at the APSA Annual Meeting, Montreal, September. Available at: https://ssrn.com/abstract=4218316

34. Oates, S., & Ramsay, G. (2024). *Seeing Red: Disinformation, Trump, and the Kremlin*. Oxford University Press.
35. Park, S., et al. (2016) 'Multimodal analysis and prediction of persuasiveness in online social multimedia', *ACM Transactions on Interactive Intelligent Systems*, 6(3), article 25.
36. Pastor-Galindo, J., et al. (2025) 'Understanding foreign information manipulation and interference: Quantitative analysis of FIMI incidents', *Computers & Security*, 140, article 103789.
37. Paul, C. & Matthews, M. (2016). The Russian "firehose of falsehood" propaganda model: Why it might work and options to counter it. *RAND Corporation Perspective Series*. Available at: https://www.rand.org/pubs/perspectives/PE198.html
38. Petty, R. E., & Cacioppo, J. T. (1986). 'The Elaboration Likelihood Model of Persuasion'. *Advances in Experimental Social Psychology*, 19, 123-205. https://doi.org/10.1016/S0065-2601(08)60214-2
39. Pew Research Center (2023) 'Americans' experiences with news on podcasts', *Pew Research Center*, 18 April. https://www.pewresearch.org/journalism/2023/04/18/americans-experience-with-news-on-podcasts/
40. Pomerantsev, P. (2024) *How to Win an Information War: The Propagandist Who Outwitted Hitler*. London: Faber & Faber.
41. Rid, T. (2020). *Active Measures: The Secret History of Disinformation and Political Warfare*. New York: Farrar, Straus and Giroux.
42. Rizwan, A., et al. (2025) 'Deciphering toxicity in political podcasts: A computational framework', *Journal of Computer-Mediated Communication*, 30(2), article zmab015.
43. Robertson, T. and Pelechaty, T. (2022) *Addressing attribution: Theorizing a model to identify Russian disinformation campaigns online*. Calgary: Canadian Global Affairs Institute. ISBN: 978-1-77397-263-3.
44. Schramm, H., et al. (2024) 'Parasocial interactions and relationships: A comprehensive meta-analysis', *Frontiers in Psychology*, 15, article 1128562.
45. SecAlliance (2024) 'Evolution of Russian influence operations 2016-2024', *SecAlliance Threat Intelligence Report*.
46. Shah, N., Willick, D., and Mago, V. (2018) 'A framework for social media data analytics using Elasticsearch and Kibana', *Wireless Networks*. doi:10.1007/s11276-018-01896-2
47. Sharon, T. and John, N. A. (2024) 'Manufacturing intimacy: Parasocial relationships in the podcast ecosystem', *New Media & Society*, 26(4), 1892-1911.
48. Shekhovtsov, A. (2017). *Russia and the Western Far Right: Tango Noir*. Routledge.
49. Sienkiewicz, M., & Jaramillo, D. L. (2019). Podcasting, the intimate self, and the public sphere. *Popular Communication*, 17(4), 268-272.
50. Solopova, A., Wang, J., and Tsvetkov, Y. (2024) 'Adaptive sarcasm detection for political communication', *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2150–2165.
51. Starbird, K., Arif, A., & Wilson, T. (2019). 'Disinformation as Collaborative Work: Surfacing the Participatory Nature of Strategic Information Operations'. *Proceedings of the ACM on Human-Computer Interaction (CSCW)*.
52. Stone, D. A. (1989). 'Causal Stories and the Formation of Policy Agendas'. *Political Science Quarterly*, 104(2), 281–300. https://doi.org/10.2307/2151585

53. Surjeet, S., Kumar, N., and Li, Z. (2025) 'Scalable fake news detection using curriculum learning and deep feature extraction', *Expert Systems with Applications*, 240, 122456.
54. Swartz, C. (2025). *DissertationSep2025: The repository to accompany my dissertation*. GitHub. https://github.com/syntax606/DissertationSep2025
55. Tucker, J. A., Guess, A., Barberá, P., Vaccari, C., Siegel, A., Sanovich, S., Stukal, D., & Nyhan, B. (2018). *Social media, political polarization, and political disinformation: A review of the scientific literature*. Menlo Park, CA: William and Flora Hewlett Foundation. doi:10.2139/ssrn.3144139
56. U.S. Department of Justice (2024) 'United States v. Kalashnikov and Afanasyeva', Criminal Indictment, Case No. 1:24-cr-00278, Eastern District of Tennessee.
57. U.S. Department of Justice. (2024) *Indictment in United States v. Tenet Media (24CRIM 519)*. Washington, D.C.: DOJ.
58. Valisa, P. (2025) 'Cultural backlash and foreign influence: Russian narratives in Italian far-right discourse', *European Journal of Communication*, 40(1), 23-47.
59. Wang, P., et al. (2021) 'Audio mining: The role of vocal tone in persuasion', *Journal of Consumer Research*, 48(2), 189-211.
60. Wirtschafter, V. (2021) 'Podcasts and the misinformation problem', *Columbia Journalism Review*, 15 November.
61. Wirtschafter, V. (2023) 'Audible reckoning: How top political podcasters spread unsubstantiated and false claims', *Brookings Institution TechStream*, February. Available at: https://www.brookings.edu/articles/audible-reckoning-how-top-political-podcasters-spread-unsubstantiated-and-false-claims/
62. Wylie, C. (2019). *Mindfck: Cambridge Analytica and the Plot to Break America*. New York: Random House.
63. Yang, J., Han, S. C., and Poon, J. (2022) 'A survey on extraction of causal relations from natural language text', *Knowledge and Information Systems*, 64, 1161–1186. doi:10.1007/s10115-022-01665-w
64. Yang, Y., McCabe, S., & Hindman, M. (2024). 'Does Russian Propaganda Lead or Follow? Topic Coverage, User Engagement, and RT and Sputnik's Agenda Influence on US Media'. *The International Journal of Press/Politics*, 0(0). https://doi.org/10.1177/19401612241271074
65. You, C., Zhu, D., Sun, Y., Ye, A., Wu, G., Cao, N., Qiu, J., and Zhou, H. M. (2019) 'SNES: Social-network-oriented public opinion monitoring platform based on ElasticSearch', *Computers, Materials & Continua*, 61(3), 1271–1283. doi:10.32604/cmc.2019.06133